

Malthus Meets Romer: Population Growth Under Information Frictions

Anson Zhou*

May 2026

[\[Click here for the latest version\]](#)

Abstract

Households disagree about whether larger populations enhance productivity (Romerians) or diminish it (Malthusians). I embed this disagreement into a Barro–Becker fertility model with global games. Wrong beliefs generate distinct pathologies: in Romerian worlds, Malthusians cause prolonged stagnation through an *informational poverty trap* in which the signals that would resolve disagreement are uninformative in the trap region; in Malthusian worlds, Romerians generate non-monotone dynamics, including boom-bust cycles. Optimal communication policy is truthful under dynastic welfare but shaded under utilitarian ones. Survey evidence confirms Malthusian beliefs correlate with lower fertility.

JEL Classification: D83, J13, O40, C72

Keywords: Fertility, heterogeneous beliefs, scale effects, global games, informational traps, endogenous population welfare

*[Anson Zhou](#): Faculty of Business and Economics, The University of Hong Kong. I thank the HKU Jockey Club Enterprise Sustainability Global Research Institute for financial support.

1 Introduction

How many children should a family have? The answer turns on what one believes about the economic consequences of population size. A parent who thinks, with [Romer \(1990\)](#), that more people generate more ideas and faster growth may view an additional child as contributing to future prosperity. A parent who fears, with [Malthus \(1798\)](#), that more people strain finite resources may view the same child as hastening collective impoverishment. Under intergenerational altruism, these beliefs translate directly into fertility choices, and the resulting population dynamics feed back into productivity outcomes that either validate or refute them. This paper studies the consequences of that feedback loop when households hold heterogeneous beliefs about the population–productivity relationship.

I formalize the idea in an overlapping generations model in which households make fertility decisions under uncertainty about a single structural parameter: the scale effect of population on productivity growth. Households observe the same aggregate data but interpret it through different priors.¹ “Romerians” believe that larger populations accelerate innovation and growth; “Malthusians” believe that larger populations depress living standards. Beliefs are transmitted across generations through a combination of cultural inheritance and Bayesian updating from observed outcomes. Following [Becker \(1960\)](#), children are costly in time rather than in goods: a parent devotes a fraction $\psi(n)$ of a unit time endowment to child-rearing and supplies the remainder to the labor market. Productivity itself evolves according to a multiplicative log-normal process whose drift depends on population, so today’s fertility shapes tomorrow’s productivity through the size of the next generation. Together, these ingredients yield a tractable framework in which beliefs, fertility, and growth co-evolve.

The model rests on a single behavioral premise: beliefs about population’s economic consequences shape fertility. Before drawing out theoretical implications, I establish this premise empirically. Using data from the General Social Survey (GSS), I document robust negative correlations between Malthusian beliefs—measured by agreement with the statement that Earth cannot continue to support population growth—and both ideal and actual fertility. Individuals with Malthusian beliefs report ideal fertility about 0.27 children lower than non-Malthusians and have 0.29 fewer children, gaps of roughly 11 and 16 percent of the respective sample means. The relationships persist after controlling for age, education, sex, and survey year, and they show up across demographic groups. Notably, the effect on

¹I treat heterogeneous priors over the structural parameter governing the population–productivity relationship as a primitive feature of the environment, not a transitional state of asymmetric information that learning will quickly erase. This stance follows recent work in macroeconomics and finance that takes survey-measured belief heterogeneity as the input to a model of long-run risks rather than as the residual of an information-friction model ([Bakkensen et al., 2025](#); [Gans, 2026](#); [Agrawal et al., 2025](#)).

actual fertility is at least as large as the effect on stated preferences, which suggests that beliefs translate into consequential behavior rather than cheap talk.²

With this empirical foundation in place, I turn to the theoretical analysis. The strategic structure of fertility partitions the state space into three regions with distinct equilibrium properties. In the *Malthusian-dominated region*, where the share of Romerians is below a threshold ϕ^* determined by the calibration, the aggregate best response is downward-sloping and the equilibrium is unconditionally unique. In the *contraction region* between ϕ^* and a second threshold ϕ^{**} , the aggregate best response satisfies $0 \leq \rho(\phi) < 1$, and uniqueness follows from a contraction argument. In the *multiplicity region* above ϕ^{**} , strategic complementarity among Romerians dominates, multiple equilibria coexist, and selection becomes a substantive question rather than a refinement of a refinement.³ The paper handles all three regions and uses a global-games refinement à la Frankel et al. (2003) to select a risk-dominant equilibrium in the multiplicity region; the existence step uses the single-crossing approach of Athey (2001).

Within this equilibrium structure, the main theoretical contribution is to show that wrong beliefs generate qualitatively different pathologies depending on which worldview is correct. The asymmetry arises from a feature of Bayesian learning about scale effects: the informativeness of productivity signals depends on population size. When population sits near the reference level $\bar{N}$, productivity growth reveals almost nothing about whether scale effects are positive or negative. Learning therefore requires population to deviate substantially from this uninformative region, but such deviations require precisely the optimistic beliefs that pessimism suppresses. This interaction between beliefs, fertility, and the informativeness of data is the engine driving the paper’s main results.

Consider first a Romerian world, where scale effects are truly positive. Here, Malthusian minorities do not merely slow convergence to efficiency—they can trap the economy in prolonged stagnation. Pessimistic beliefs suppress fertility, keeping population near $\bar{N}$. A population near $\bar{N}$ generates uninformative signals about the population–productivity relationship, and uninformative signals fail to correct pessimistic beliefs. The trap is self-reinforcing: the economy remains stuck not because Malthusianism is correct, but because population never deviates enough from $\bar{N}$ to reveal that it is wrong. I call this the *informational poverty trap*. While correct beliefs eventually prevail almost surely, the expected

²The use of survey-based belief proxies follows the methodology of Howe et al. (2015) as adapted in Bernstein et al. (2019), Baldauf et al. (2020), and Bakkensen and Barrage (2022); limitations of the correlational approach are discussed in Section 3.

³Whether the multiplicity region is empirically active depends on the strength of scale effects. Under Jones-style semi-endogenous growth calibrations (Jones, 1995; Bloom et al., 2020), the multiplicity region is empty; under Kremer-style or Galor–Weil-style calibrations (Kremer, 1993; Galor, 2011), it is occupied at high ϕ .

escape time can exceed any relevant historical horizon.

Now consider the mirror case: a Malthusian world, where scale effects are truly negative. Here the dynamics differ starkly. Demographic selection always favors optimists, who have more children regardless of whether their optimism is warranted. Bayesian learning, by contrast, favors pessimists when population growth leads to poor outcomes. The two forces pull the Romerian share ϕ_t in opposite directions, and the balance between them itself depends on the state. The paper’s second main theoretical result formalizes this opposition: in any Malthusian world, the expected drift in ϕ_t cannot be uniformly signed, because the learning channel vanishes at $N_t = \bar{N}$ (where selection dominates) and grows quadratically in $\ln(N_t/\bar{N})$ thereafter. Monotone convergence is therefore ruled out, and the system inherits the potential for non-monotone dynamics, including the boom-bust cycles familiar from pre-industrial demography. During expansions, optimists outbreed pessimists and population grows. Eventually the economy hits constraints, outcomes deteriorate, learning corrects beliefs toward Malthusianism, and fertility falls. Recovery begins, and the pattern can repeat.

Taken together, the two cases reveal a sharp duality: *in Romerian worlds, Malthusians cause stagnation; in Malthusian worlds, Romerians generate non-monotone dynamics, including cycles.* Wrong beliefs are always costly, but the nature of the cost depends on the direction of the error.

This duality is more than a theoretical curiosity. It offers a unified lens for interpreting demographic history—though one that admits two competing readings of the pre-industrial Malthusian era. The conventional reading is that the pre-industrial world was genuinely Malthusian: scale effects were negative, and the cycles documented by economic historians (Clark, 2007; Voigtländer and Voth, 2013) reflect the boom-bust dynamics of a true Malthusian regime. On this view, the Industrial Revolution marked a genuine technological regime change. The alternative reading, made possible by the informational poverty trap, is that the world was always Romerian, but humanity was stuck in a trap where population never deviated enough from the reference to generate the information needed to escape pessimism. On this view, the Industrial Revolution was not a technological regime change but an *informational escape*: some combination of circumstances pushed population into the informative region, beliefs updated toward Romerianism, and the trap dissolved. Both readings are consistent with the historical evidence, and distinguishing them carries profound implications for understanding the modern fertility decline—which the model can interpret as a partial drift back toward the trap, an efficient response to a renewed Malthusian regime, or a quantity-quality substitution unrelated to scale effects.

The theoretical results naturally raise a policy question: if wrong beliefs can trap economies

in stagnation or generate damaging cycles, can public communication help? Answering this question requires first confronting what “help” even means when the population itself is endogenous to the policy. As [Goloso et al. \(2007\)](#) emphasize, standard welfare criteria become problematic when the set of agents depends on policy.⁴ I analyze optimal policy under three welfare criteria that span a substantively distinct normative spectrum: a dynastic criterion that respects Barro–Becker preferences, a total utilitarian criterion that aggregates utility across all individuals who exist under the true law of motion, and an average utilitarian criterion. The contrast suffices to bring out genuinely distinct policy implications. A dynastic planner is fundamentally an information provider; a total utilitarian planner is an externality corrector. The former is indifferent to belief correction beyond information provision; the latter has a stake in both truth and population. An average-utilitarian benchmark sits between them.

Optimal announcement policy differs substantially across welfare criteria. A dynastic planner favors truthful communication: under subjective preferences, any informative signal is welfare-improving and there is no externality to correct. A total utilitarian planner shades announcements toward optimism (in a Romerian world) by an amount that depends on the size of the externality, the deviation of population from the reference level, and the dispersion of priors. An average-utilitarian planner shades less than the total utilitarian but more than the dynastic, because it internalizes the per-capita scale-effect spillover but not the population-size motive. The contrast across criteria is itself a substantive result: normative conclusions about the costs of wrong beliefs depend critically on which welfare standard one adopts.

The welfare effects of public announcements also depend on where the economy sits in the state space. In the contraction region, where equilibrium is unique, announcements act purely as information provision, shifting households’ posteriors toward truth (or, by design, away from it when correcting an externality). In the multiplicity region, announcements additionally play a selection-shifting role, coordinating beliefs to move the economy between low-fertility and high-fertility equilibria.

A key finding is that public information is most valuable precisely when the planner knows least: in the informational poverty trap, a credible announcement can break self-reinforcing pessimism and push population into the informative region, yet the same uninformative signals that trap households also limit what the planner can learn. This contrasts with standard results in coordination games, where public information can be harmful by inducing excessive coordination on incorrect signals ([Morris and Shin, 2002](#)). Here, the heteroge-

⁴The axiomatic foundations of population ethics, including the role of critical-level utilitarianism in addressing the “Repugnant Conclusion” of [Parfit \(1984\)](#), are developed in [Blackorby et al. \(2005\)](#).

neous strategic interaction—Romerians exhibiting complementarity, Malthusians exhibiting substitutability—provides a natural stabilizer. In the trap, even moderately noisy announcements improve welfare; in Malthusian worlds with non-monotone dynamics, higher precision is needed to avoid amplifying volatility.

Related Literature

This paper connects several literatures: endogenous fertility, heterogeneous beliefs, global games, cultural transmission, and recent work treating prior heterogeneity as a primitive of long-run economic environments.

Endogenous fertility and population scale effects. The foundational work includes [Becker \(1960\)](#), who introduced the time-cost-of-children formulation I adopt, and [Barro and Becker \(1989\)](#), who embedded fertility choice in a dynastic optimization framework. [de la Croix and Doepke \(2003\)](#) pin down the time cost of one child at approximately fifteen percent of a parent’s time endowment, which I use to calibrate the curvature parameter. Subsequent work enriches this framework with human capital ([Galor and Weil, 2000](#)), mortality transitions ([Kalemli-Ozcan, 2003](#)), and intergenerational transfers ([Lee, 2003](#)); the unified growth theory of [Galor \(2011\)](#) synthesizes these elements. On scale effects, [Romer \(1990\)](#), [Kremer \(1993\)](#), [Jones \(1995, 2022\)](#), and [Bloom et al. \(2020\)](#) provide alternative frameworks; [Ashraf and Galor \(2013\)](#) finds positive long-horizon scale effects, while [Voigtländer and Voth \(2013\)](#) documents pre-industrial Malthusian dynamics. My contribution introduces belief heterogeneity as a driver of fertility differentials and population dynamics.

Heterogeneous beliefs. [Angeletos and La’O \(2013\)](#); [Angeletos et al. \(2018\)](#), [Bordalo et al. \(2018\)](#), and [Malmendier and Nagel \(2011\)](#) examine dispersed information and expectation formation in macroeconomics; [Simsek \(2013\)](#), [Geanakoplos \(2010\)](#), and [Scheinkman and Xiong \(2003\)](#) study disagreement in finance. A more recent line of work treats heterogeneous priors over slow-moving structural parameters as a primitive: [Bakkensen et al. \(2025\)](#) embed climate-belief disagreement into mortgage origination, and [Gans \(2026\)](#) formalizes a strategic entry game in which prior heterogeneity is primitive, building on [Agrawal et al. \(2025\)](#) and [Felin and Zenger \(2009, 2017\)](#). My paper applies this approach to fertility.

Global games and cultural transmission. The global games literature ([Carlsson and Van Damme, 1993](#); [Morris and Shin, 1998, 2003](#); [Frankel et al., 2003](#)) provides tools for analyzing games with strategic complementarities under incomplete information. I adapt these

techniques to a setting that mixes complementarities (among Romerians) and substitutabilities (among Malthusians); the existence step uses [Athey \(2001\)](#)’s single-crossing argument for monotone equilibria in supermodular games with continuous actions. On cultural transmission, I draw on [Cavalli-Sforza and Feldman \(1981\)](#), [Bisin and Verdier \(2001, 2011\)](#), and [Doepke and Zilibotti \(2019\)](#).

Public information and welfare with endogenous population. [Morris and Shin \(2002\)](#) shows that public information can be harmful in coordination games by crowding out private information. [Golosov et al. \(2007\)](#) provide a systematic analysis of efficiency concepts when fertility is endogenous, showing standard welfare theorems fail under total or average utilitarianism. I build on these insights. The poverty-trap literature ([Azariadis and Drazen, 1990](#); [Banerjee and Newman, 1993](#)) studies traps from technological nonconvexities or credit constraints; the informational poverty trap I identify arises instead from endogenous uninformativeness of signals when the economy remains in a particular region of the state space.

The paper proceeds as follows. Section 2 presents the model. Section 3 provides empirical evidence. Section 4 analyzes strategic interaction and equilibrium selection. Section 5 examines long-run dynamics. Section 6 discusses demographic history. Section 7 analyzes optimal communication policy. Section 8 concludes. All proofs are in the Appendix.

2 Model

I develop a discrete-time overlapping generations model with dynastic altruism. The economy is populated by a continuum of households who make fertility decisions and differ in their beliefs about the structural relationship between population and productivity growth.

2.1 Environment

Time is discrete and indexed by $t = 0, 1, 2, \dots$. Each period, a generation of adults makes economic decisions; their children become adults in the following period.

Time Endowment and Cost of Children. Each adult is endowed with one unit of time. An adult who chooses fertility n_{it} spends a fraction $\psi(n_{it}) \in [0, 1]$ of that time on child-rearing and supplies the remainder $1 - \psi(n_{it})$ to the labor market. The implicit price of an additional child is the forgone labor income. This time-cost specification, originally from [Becker \(1960\)](#), makes the resource constraint multiplicative rather than additive, which plays

a central analytical role by making the household value function tractable to a first-order log-linear approximation.

Production. A single final good is produced using effective labor:

$$Y_t = A_t L_t, \quad L_t = N_t \cdot \int_0^1 (1 - \psi(n_{it})) di, \quad (1)$$

where A_t is total factor productivity, N_t is adult population, and L_t is aggregate effective labor supply.⁵ The competitive wage equals A_t , so a type- j household with fertility $n_{j,t}$ earns $A_t(1 - \psi(n_{j,t}))$. With no other income or expenditure, consumption is

$$c_{j,t} = A_t(1 - \psi(n_{j,t})). \quad (2)$$

TFP Dynamics. TFP evolves according to a multiplicative log-normal process,

$$A_{t+1} = A_t \cdot \exp(\bar{g} + \theta^* \ln(N_t/\bar{N}) + \eta_{t+1}), \quad \eta_{t+1} \stackrel{\text{i.i.d.}}{\sim} \mathcal{N}(0, \sigma_\eta^2), \quad (3)$$

where $\bar{g} \geq 0$ is baseline log-productivity growth, $\theta^* \in \mathbb{R}$ is the *scale-effect parameter*, $\bar{N} > 0$ is an empirically anchored *reference population*, and η_{t+1} is i.i.d. structural noise. Taking logs gives the exact identity

$$\ln A_{t+1} = \ln A_t + \bar{g} + \theta^* \ln(N_t/\bar{N}) + \eta_{t+1}. \quad (4)$$

The technology side contributes *no* Taylor remainder; any approximation in what follows enters only through the optimal fertility policy.⁶

When $\theta^* > 0$, larger populations generate faster productivity growth (Romer); when $\theta^* < 0$, larger populations slow productivity growth (Malthusian dynamic congestion). When $N_t = \bar{N}$, the deterministic component of growth equals $\bar{g}$ regardless of θ^* . Observations near $N_t = \bar{N}$ are therefore *uninformative* about the scale-effect parameter, even though they remain noisy through η_{t+1} . This feature is central to the informational poverty trap.⁷

⁵I interpret A_t broadly as “effective” TFP that incorporates any potential dilution effects from fixed factors.

⁶The reference population $\bar{N}$ appears explicitly, so the model is invariant under rescaling of the population unit. I write $g(N_t; \theta) \equiv \bar{g} + \theta \ln(N_t/\bar{N})$ for the productivity-growth function.

⁷The parameter $\theta < 0$ here represents *dynamic* congestion or environmental drag rather than classical Malthusian fixed-factor dilution operating on the level of output. Capturing the latter would require an additional production margin (a fixed factor), which is straightforward but distinct from the channel modeled here.

Households and Beliefs. The economy contains a unit mass of dynastic households indexed by $i \in [0, 1]$. Each household holds a belief θ_i about θ^* .

Assumption 1 (Belief Types). *A fraction $\phi_t \in [0, 1]$ of households are Romerians with belief $\theta_R > 0$; the remaining $1 - \phi_t$ are Malthusians with belief $\theta_M < 0$.*

A “Romerian world” has $\theta^* > 0$; a “Malthusian world” has $\theta^* < 0$. Households do not know θ^* with certainty but hold a common prior $\pi_t \in [0, 1]$ that the world is Romerian and update via Bayes’ rule. The distribution ϕ_t is common knowledge in each period; individual types are private information.

Time-Cost Function.

Assumption 2 (Time Cost of Children). *The time-cost function $\psi : \mathbb{R}_+ \rightarrow [0, 1]$ satisfies (i) $\psi(0) = 0$; (ii) $\psi'(n) > 0$ and $\psi''(n) \geq 0$; (iii) $\psi(n) < 1$ for finite n with $\lim_{n \rightarrow \infty} \psi(n) = 1$; and (iv) the marginal time cost $\psi'(n)/(1 - \psi(n))$ is strictly increasing with $\lim_{n \rightarrow \infty} \psi'(n)/(1 - \psi(n)) = \infty$.*

For tractability, I use the exponential–quadratic specification

$$\psi(n) = 1 - e^{-\frac{\gamma}{2}n^2}, \quad \gamma > 0, \quad (5)$$

which yields $\psi'(n)/(1 - \psi(n)) = \gamma n$. This identity exactly mirrors the linear marginal-cost structure of a goods-cost specification while keeping consumption bounded between zero and A_t . Following [de la Croix and Doepke \(2003\)](#), the time cost of one child is approximately fifteen percent of a parent’s time endowment, which under (5) pins down $\gamma = -2 \ln(0.85) \approx 0.325$.

Preferences. Households have dynastic preferences:

$$U_{it} = u(c_{it}) + \beta n_{it} \mathbb{E}_{it} [V_{i,t+1}], \quad (6)$$

where $u(c) = \ln c$ and $\beta \in (0, 1)$ is the dynastic discount factor. Throughout, I assume equilibrium fertility satisfies $\beta n_j^* < 1$. I set $\beta = 0.30$, interpreted as a structural per-generation altruism parameter.⁸ Substituting $c_{it} = A_t(1 - \psi(n_{it}))$ gives

$$\ln c_{it} = \ln A_t + \ln(1 - \psi(n_{it})). \quad (7)$$

⁸This value sits at the lower end of the range used in the quantitative Barro–Becker literature ([Manuelli and Seshadri, 2009](#); [Lee and Seshadri, 2019](#); [Daruich, 2018](#)). It is conservative for the present analysis: a higher altruism parameter would amplify the strategic complementarity slope d_R in (18) and *enlarge* the multiplicity region.

The first term depends only on aggregate productivity, the second only on own fertility through the time cost.

Population Dynamics.

$$N_{t+1} = \bar{n}_t \cdot N_t, \quad \bar{n}_t = \phi_t n_{R,t} + (1 - \phi_t) n_{M,t}, \quad (8)$$

where $n_{R,t}$ and $n_{M,t}$ denote Romerian and Malthusian fertility.

2.2 Belief Formation and Intergenerational Transmission

Cultural Transmission.

Assumption 3 (Cultural Transmission). *A child born to a type- j household adopts type j with probability $\lambda \in [0, 1]$ and forms beliefs through learning with probability $1 - \lambda$.*

Bayesian Learning. Children who learn form beliefs by Bayes' rule. Observed log-productivity growth is

$$s_t \equiv \ln(A_{t+1}/A_t) = \bar{g} + \theta^* \ln(N_t/\bar{N}) + \eta_{t+1}. \quad (9)$$

Given s_t and N_t , the likelihood under belief θ is

$$\ell(s_t \mid \theta, N_t) = \frac{1}{\sqrt{2\pi}\sigma_\eta} \exp\left(-\frac{(s_t - \bar{g} - \theta \ln(N_t/\bar{N}))^2}{2\sigma_\eta^2}\right), \quad (10)$$

and Bayes' rule gives

$$\pi_{t+1} = \frac{\pi_t \cdot \ell(s_t \mid \theta_R, N_t)}{\pi_t \cdot \ell(s_t \mid \theta_R, N_t) + (1 - \pi_t) \cdot \ell(s_t \mid \theta_M, N_t)}. \quad (11)$$

Children who learn adopt Romerian beliefs with probability π_{t+1} and Malthusian beliefs with probability $1 - \pi_{t+1}$.

Signal Informativeness. The Kullback–Leibler divergence between the signal distributions under the two beliefs is

$$\mathcal{I}_t \equiv \text{KL}(\ell(\cdot \mid \theta_R, N_t) \parallel \ell(\cdot \mid \theta_M, N_t)) = \frac{(\theta_R - \theta_M)^2 [\ln(N_t/\bar{N})]^2}{2\sigma_\eta^2}. \quad (12)$$

When $N_t = \bar{N}$, $\mathcal{I}_t = 0$: the signal is completely uninformative about θ^* . As $|\ln(N_t/\bar{N})|$ grows, the signal becomes increasingly informative. This dependence is the source of the

informational poverty trap and, in Malthusian worlds, of the opposing forces analyzed in Section 5.

Evolution of Belief Distribution. Combining cultural transmission and learning, the Romerian share evolves as

$$\phi_{t+1} = \lambda \phi_t^{birth} + (1 - \lambda) \pi_{t+1}, \quad (13)$$

where $\phi_t^{birth} = \phi_t n_{R,t} / \bar{n}_t$ is the fraction of children born to Romerians.

2.3 Independence of Priors

A modeling choice deserves explicit mention. In the model above, a household's prior θ_i and other households' priors are drawn independently from the population distribution induced by (ϕ_t, π_t) . Knowing one's own θ_i does not lead a household to update its belief about other households' priors.

This assumption matches a world in which prior heterogeneity comes from idiosyncratic sources: different parental theories, different tacit knowledge, different analogical reasoning, and different tastes for risk-taking in the construction of a mental model. A household that happens to be very optimistic learns nothing about other households' optimism from its own conviction, because its conviction is generated by an idiosyncratic internal theory rather than a common external signal.⁹ The cultural-transmission parameter λ introduces a within-dynasty correlation that does not affect the cross-section independence assumption. The public-validation extension in Section 7 introduces a different limited form of correlation: a public announcement shifts every household's prior by the same amount while leaving the residual idiosyncratic component independent.

2.4 Equilibrium

Definition 1 (Equilibrium). *An equilibrium consists of fertility policies $n_R(A_t, N_t, \phi_t, \pi_t)$, $n_M(A_t, N_t, \phi_t, \pi_t)$; aggregate fertility $\bar{n}(A_t, N_t, \phi_t, \pi_t)$; and laws of motion for $(A_{t+1}, N_{t+1}, \phi_{t+1}, \pi_{t+1})$ given by (3), (8), (13), and (11), such that (a) each household's fertility maximizes (6) subject to (2) given subjective beliefs and the aggregate fertility function; (b) expectations about aggregate fertility are consistent with the equilibrium aggregate; and (c) actual laws of motion use the true θ^* and realized η_{t+1} .*

⁹This is the natural counterpart of the *theory-based* view of strategic disagreement (Felin and Zenger, 2009, 2017). If, by contrast, prior heterogeneity were generated by a common source—a publicly observable disclosure—then a household's high own prior would rationally shift its belief about others' priors upward, and individual priors would be positively correlated conditional on (ϕ_t, π_t) . The clean closed-form best-response slopes derived in Section 4 would no longer obtain in the same form.

The Bellman below treats a type- j parent's continuation value as the value function of a type- j child, even though the child may switch with probability $1 - \lambda$.¹⁰

2.5 Household Optimization

The Bellman for a type- j household reads

$$V_j(A, N, \phi, \pi) = \max_{n \geq 0} \{ \ln A + \ln(1 - \psi(n)) + \beta n \cdot \mathbb{E}_j[V_j(A', N', \phi', \pi')] \}. \quad (14)$$

I work with a first-order log-linear approximation around a balanced-growth reference. I conjecture

$$V_j(A, N, \phi, \pi) \approx v_j(\phi, \pi) + \xi_{A,j} \ln A + \xi_{N,j} \ln(N/\bar{N}), \quad (15)$$

with $(\xi_{A,j}, \xi_{N,j})$ and $v_j(\phi, \pi)$ determined by coefficient matching. The technology side is exactly linear in $(\ln A, \ln(N/\bar{N}))$, so any deviation from (15) comes only from the dependence of optimal fertility n_j^* on the state.

The first-order condition for fertility is

$$\frac{\psi'(n_j)}{1 - \psi(n_j)} = \beta \mathbb{E}_j[V_j(A', N', \phi', \pi')], \quad (16)$$

which under (5) simplifies to $\gamma n_j = \beta \mathbb{E}_j[V_j]$.

Lemma 1 (Equilibrium Fertility, Local Coefficients). *Under Assumptions 1–3 and the local log-linear conjecture (15), the value function coefficients evaluated at the balanced-growth reference are*

$$\xi_{A,j} = \frac{1}{1 - \beta n_j^*}, \quad \xi_{N,j} = \frac{\beta n_j^* \theta_j}{(1 - \beta n_j^*)^2}. \quad (17)$$

Equilibrium fertility satisfies a fixed-point equation in n_j^ and can be written as $n_j^* = n^* h(\theta_j; \bar{n})$ with h increasing in θ . Moreover, $n_R^* > n_M^*$ for all equilibrium values of $\bar{n}$.*

The proof is in Appendix A. Households who believe population enhances productivity (Romerians) perceive a higher return to children and choose higher fertility. This testable prediction motivates the empirical analysis in the next section.

¹⁰An exact formulation would replace $\mathbb{E}_j[V_j]$ by $\lambda \mathbb{E}_j[V_j] + (1 - \lambda) \mathbb{E}_j[V_{learn}]$, where V_{learn} is a type-mixture continuation. The two formulations coincide when λ is close to 1 or when V_R and V_M are close in the relevant region; both conditions hold near the steady states studied below.

3 Empirical Evidence: Beliefs and Fertility

A central prediction of the model is that Malthusian beliefs reduce fertility. This section examines empirical support using survey data.

3.1 Data

I use the General Social Survey (GSS), a nationally representative survey of U.S. adults. The key belief variable comes from the question: “The earth cannot continue to support population growth at its present rate.” Respondents indicate agreement on a five-point scale, and I classify those who agree or strongly agree as holding *Malthusian beliefs*. This question was asked in the 2000 and 2010 waves.¹¹ I examine two fertility measures: stated preferences (“ideal number of children”) and realized fertility (actual number of children). The sample contains 1,822 observations.

3.2 Descriptive Patterns

Table 1 presents summary statistics. Mean ideal fertility is 2.50 (s.d. 0.87) and mean actual fertility is 1.81 (s.d. 1.64). Approximately 55 percent of respondents hold Malthusian beliefs.

Table 1. Summary Statistics

Variable	N	Mean	Std. Dev.	Min	Max
Ideal number of children	1,822	2.50	0.87	0	7
Actual number of children	1,822	1.81	1.64	0	8
Malthusian belief (binary)	1,822	0.55	0.50	0	1
Malthusian belief (1–5 scale)	1,822	3.44	1.03	1	5
Age	1,822	46.8	17.3	18	89
Years of education	1,822	13.3	2.96	0	20
Female	1,822	0.56	0.50	0	1

Notes: Data from GSS, 2000 and 2010 waves. Malthusian belief (binary) equals one if respondent agrees or strongly agrees that Earth cannot continue to support population growth.

Figure 1 displays mean ideal and actual fertility by Malthusian belief intensity. There is a clear monotonic negative relationship between belief strength and both fertility measures. Respondents who strongly disagree report an average ideal of 2.98 children and have 1.76 children, while those who strongly agree report 2.24 and 1.52.

¹¹The use of survey-measured belief proxies has a substantial track record in environmental and climate finance (Howe et al., 2015; Bernstein et al., 2019; Baldauf et al., 2020; Bakkensen and Barrage, 2022; Bakkensen et al., 2025), with validation evidence that aggregated belief measures correlate strongly with door-to-door




Figure 1. Fertility by Malthusian Belief Intensity

Notes: Data from GSS, 2000 and 2010 waves (N = 1,822). Bars show mean ideal (navy) and actual (maroon) number of children for each level of the Malthusian belief scale. Error bars indicate 95% confidence intervals.

A natural concern is that this relationship is driven by education. Figure 2 shows mean ideal fertility by education group and belief type. Within every education category, Malthusians report lower ideal fertility than non-Malthusians, with gaps ranging from 0.20 to 0.45 children.

3.3 Regression Analysis

Table 2 presents linear regressions of both ideal and actual fertility on Malthusian beliefs with demographic controls.¹²

Malthusians report ideal fertility 0.27 children lower than non-Malthusians (column 1), about 11 percent of the sample mean. The coefficient on actual fertility is -0.29 (column 3), implying Malthusians have approximately 16 percent fewer children. The coefficient on actual fertility is *larger* in absolute value than the coefficient on ideal fertility, suggesting individual surveys.

¹²All specifications include both age and age squared, so the coefficient on age in the table reflects the linear component only. The marginal effect of age combines linear and quadratic components and is therefore age-specific.




Figure 2. Ideal Number of Children by Education and Malthusian Belief

Notes: Bars show mean ideal number of children for non-Malthusians (navy) and Malthusians (maroon) within each education category.

beliefs translate into behavior with at least as much force as they shape stated preferences.

3.4 Completed Fertility and Heterogeneity

A concern with actual fertility is that younger respondents may not have completed child-bearing. Table 3 restricts the sample to women aged 45+ and men aged 50+. Among those with completed fertility, Malthusians have 0.26 fewer children than non-Malthusians, similar to the full-sample estimate.

Table 4 examines heterogeneity by education. The belief-fertility relationship is present across education groups for both outcome measures. For ideal fertility, the coefficient is somewhat larger for less educated respondents; for actual fertility, the pattern reverses, with the coefficient larger among the more educated. The descriptive pattern is consistent with the interpretation that Malthusian beliefs affect fertility preferences across the education distribution but translate into behavior more strongly among the more educated, perhaps because they face fewer constraints on achieving their fertility targets.

Table 2. Effect of Malthusian Beliefs on Fertility

	Ideal Fertility		Actual Fertility	
	(1)	(2)	(3)	(4)
Malthusian	−0.268*** (0.041)		−0.292*** (0.066)	
Malthusian (continuous)		−0.153*** (0.021)		−0.146*** (0.033)
Age	−0.014** (0.006)	−0.015** (0.006)	0.112*** (0.010)	0.111*** (0.010)
Female	0.015 (0.040)	0.026 (0.040)	0.091 (0.068)	0.104 (0.068)
Education	−0.047*** (0.007)	−0.046*** (0.007)	−0.144*** (0.013)	−0.143*** (0.013)
Year FE	Yes	Yes	Yes	Yes
Observations	1,822	1,822	1,822	1,822
R^2	0.060	0.069	0.262	0.262

Notes: All models include age squared and year fixed effects. Robust standard errors in parentheses. The reported age coefficient is the linear term only. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

Table 3. Actual Fertility by Completed Fertility Status

	Full Sample (1)	Completed Fertility (2)	Not Yet Completed (3)
Malthusian	−0.292*** (0.066)	−0.257** (0.109)	−0.288*** (0.079)
Observations	1,822	848	974
R^2	0.262	0.112	0.266

Notes: Dependent variable is actual number of children. All models control for age, age squared, sex, education, and year. Robust standard errors in parentheses. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

3.5 Interpretation and Limitations

The empirical findings are consistent with the model’s central prediction: individuals who believe population growth is unsustainable both desire and have fewer children. The relationships are economically meaningful and statistically robust. Two features deserve emphasis. First, the relationship holds for *actual* fertility, not just stated preferences. Second, the coefficient on actual fertility is at least as large as the coefficient on ideal fertility.

Several limitations are worth flagging. The evidence remains correlational and could

Table 4. Heterogeneity in the Belief–Fertility Relationship

	Ideal Fertility		Actual Fertility	
	HS or Less	Some College+	HS or Less	Some College+
Malthusian	−0.318*** (0.069)	−0.235*** (0.049)	−0.209* (0.115)	−0.392*** (0.081)
Observations	821	1,001	821	1,001
R^2	0.045	0.029	0.163	0.240

Notes: Dependent variable is ideal (columns 1–2) or actual (columns 3–4) number of children. Robust standard errors in parentheses. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

reflect reverse causality or omitted-variable bias. The cross-sectional nature of the data precludes tracking how belief changes affect subsequent fertility decisions. The belief measure also captures concern about global population sustainability, not specifically the scale-effect parameter θ .¹³ Despite these limitations, the evidence provides suggestive support for the belief channel.

4 Strategic Interaction and Equilibrium Selection

This section analyzes strategic interaction in fertility choices and partitions the equilibrium space into three economically distinct regions. Three results organize the rest of the paper. First, Romerians exhibit strategic complementarity while Malthusians exhibit strategic substitutability, so the slope of the aggregate best response is a continuous function of the Romerian share that switches sign at an interior threshold. Second, two thresholds ϕ^* and ϕ^{**} partition the share space into a Malthusian-dominated region of unconditional uniqueness, a contraction region of conditional uniqueness, and a Romerian-dominated multiplicity region whose presence depends on the strength of scale effects. Third, in the multiplicity region a global-games refinement à la Frankel et al. (2003) selects a risk-dominant equilibrium.

4.1 The Fertility Game

Consider fertility choice as a game among households. Each household i chooses n_i to maximize expected utility, and aggregate fertility is $\bar{n} = \int n_i di$. The household’s payoff

¹³A natural template for stronger identification would combine survey-based belief proxies with high-resolution geographic variation in exposure to belief-relevant shocks (extreme weather events, salient environmental news, or pronatalist policy changes), or use repeated cross-sections to track within-county changes in beliefs against subsequent fertility. The GSS data alone do not support such designs, and I view causal identification as a priority for follow-up empirical work.

depends on $(n_i, \bar{n})$ through a single channel: higher $\bar{n}$ raises $N' = \bar{n}N$, which through (3) affects A'' and subsequent productivity. Let $U_j(n_j, \bar{n})$ denote the expected utility of a type- j household when aggregate fertility is $\bar{n}$. Type j exhibits *strategic complementarity* if $\partial^2 U_j / \partial n_j \partial \bar{n} > 0$ and *strategic substitutability* if the cross-partial is negative.

4.2 Best Responses and the Sign of Strategic Interaction

Let $n_j(\bar{n}^e)$ denote the fertility choice of type j when expected aggregate fertility is $\bar{n}^e$. Rational expectations require $\bar{n}^e = \bar{n}$. Under the local log-linear approximation, the marginal effect of expected aggregate fertility on the continuation value is $\partial \mathbb{E}_j[V_j] / \partial \bar{n}^e = \xi_{N,j} / \bar{n}^e$. From (17), the sign of $\xi_{N,j}$ matches the sign of θ_j : $\xi_{N,M} < 0$ for Malthusians and $\xi_{N,R} > 0$ for Romerians. Since the household FOC has a strictly increasing left-hand side in n_j , the sign of $dn_j^* / d\bar{n}^e$ coincides with the sign of $\xi_{N,j}$.

Proposition 1 (Strategic Interaction). *Under Assumptions 1–3: (i) Romerians exhibit strategic complementarity, $dn_R^* / d\bar{n}^e > 0$; (ii) Malthusians exhibit strategic substitutability, $dn_M^* / d\bar{n}^e < 0$. The closed-form local slopes are*

$$d_j \equiv \left. \frac{dn_j^*}{d\bar{n}^e} \right|_{ref} = \frac{\beta^2 n_j^* \theta_j}{\gamma \bar{n} (1 - \beta n_j^*)^2}. \quad (18)$$

The proof is in Appendix A. A Romerian who expects higher aggregate fertility expects higher future population, which under Romerian beliefs implies faster productivity growth and a higher continuation value of each child. A Malthusian who expects higher aggregate fertility expects slower productivity growth, which lowers the continuation value and reduces fertility.

4.3 Three Equilibrium Regions

Define the aggregate best-response slope

$$\rho(\phi) \equiv \phi d_R + (1 - \phi) d_M = \phi d_R - (1 - \phi) |d_M|, \quad (19)$$

with $d_R > 0 > d_M$. $\rho(\phi)$ crosses zero at

$$\phi^* = \frac{|d_M|}{d_R + |d_M|}, \quad (20)$$

and crosses one at

$$\phi^{**} = \frac{1 + |d_M|}{d_R + |d_M|}, \quad (21)$$

provided the right-hand side is below one—which holds iff $d_R > 1$. The two thresholds carve the unit interval into three economically distinct regions.

Theorem 1 (Equilibrium Existence and Uniqueness by Region). *Suppose Assumptions 1–3 hold and Proposition 1 applies at the relevant reference. Then:*

- (i) **Malthusian-dominated region**, $\phi < \phi^*$. *The aggregate best response $BR(\bar{n}^e) = \phi n_R^*(\bar{n}^e) + (1 - \phi)n_M^*(\bar{n}^e)$ is strictly decreasing in a neighborhood of any candidate fixed point. The equilibrium is locally unique unconditionally on the magnitude of $\rho(\phi)$.*
- (ii) **Contraction region**, $\phi^* \leq \phi \leq \phi^{**}$ (or all $\phi \geq \phi^*$ when $d_R \leq 1$). $0 \leq \rho(\phi) < 1$, *the contraction-mapping argument applies, and the equilibrium is locally unique.*
- (iii) **Multiplicity region**, $\phi > \phi^{**}$ (only when $d_R > 1$). $\rho(\phi) > 1$ and *BR can intersect the 45-degree line transversally at multiple points, generating an odd number of locally ordered equilibria of which the extreme two are stable.*

The proof is in Appendix A. In the Malthusian-dominated region, optimist coordination is too thin to generate the strategic complementarity needed for multiple stable points. In the contraction region, complementarity is bounded, so the standard contraction argument delivers uniqueness. Only in the Romerian-dominated region with sufficiently strong scale effects does complementarity dominate enough to generate multiplicity.

When does the multiplicity region exist? Theorem 1(iii) requires $d_R > 1$, which substituting (18) reads

$$\theta_R > \theta_R^{crit} \equiv \frac{\gamma \bar{n} (1 - \beta n_R^*)^2}{\beta^2 n_R^*}. \quad (22)$$

At the calibration $(\beta, \gamma) = (0.30, 0.325)$ with $n_R^* = 1.2$ and $\bar{n} \approx 1$, $\theta_R^{crit} \approx 1.23$. Under standard semi-endogenous-growth calibrations the multiplicity region is empty, while under Kremer–Galor–Weil calibrations it is occupied at high ϕ . Both cases are interesting.

4.4 Stability and Global-Games Selection

Within the multiplicity region, local-stability provides a first-pass refinement: an equilibrium is stable iff $|\rho(\phi)| < 1$ at the slopes evaluated locally at that equilibrium. With $\rho(\phi) > 1$ on average, the typical configuration is two stable equilibria flanking an unstable middle equilibrium. Stability eliminates the middle equilibrium but leaves the two stable selections to be coordinated upon.

To break this remaining indeterminacy, I adopt the global-games refinement of Frankel et al. (2003). The key technical step is a reduction: because Malthusians have a strictly

downward-sloping best response and a unique fixed point given any aggregate fertility level, their best response can be substituted out. The reduced game among Romerians is super-modular and falls within the scope of Frankel et al. (2003).¹⁴

Theorem 2 (Global-Games Selection in the Multiplicity Region). *Consider the multiplicity region of Theorem 1(iii). Suppose households observe private signals $\phi_i = \phi + \zeta_i$ with $\zeta_i \sim \mathcal{N}(0, \sigma_\zeta^2)$ and the prior on ϕ is diffuse. As $\sigma_\zeta \rightarrow 0$, iterated deletion of strictly dominated strategies in the reduced Romerian subgame selects a unique threshold equilibrium with cutoff $\phi^{rd}(\theta_R, \theta_M, \beta, \gamma) \in (\phi^{**}, 1)$ at which Romerians coordinate on the high-fertility selection above the cutoff and on the low-fertility selection below. The selected equilibrium is risk-dominant in the sense of Frankel et al. (2003).*

The proof is in Appendix A. The result restores a unique equilibrium selection across the entire ϕ -space.

Why selection matters for policy. Outside the multiplicity region, an announcement plays a single role: it shifts households' subjective posteriors over θ , and through them the unique equilibrium fertility. Inside the multiplicity region, an announcement plays a second role: by shifting the prior over ϕ , it can move the economy across the cutoff ϕ^{rd} and switch the selection. The selection-shifting role is most relevant when the prior ϕ is near ϕ^{rd} and absent when the prior is far from ϕ^{rd} . Section 7 treats both roles.

4.5 Map of Regions and Empirical Relevance

Table 5 summarizes the partition. The trap region of Section 5 typically features ϕ below the critical share at which aggregate fertility just reaches replacement, which is well below ϕ^{**} for any standard calibration; the trap therefore sits inside the contraction (or Malthusian-dominated) region, and the policy intuition for the trap is unaffected by multiplicity. The interior steady state of the Malthusian world also sits in the contraction region for standard calibrations.

The stabilizing role of Malthusians. Because Malthusian best responses slope downward, the Malthusian fraction $1 - \phi$ acts as a continuous shifter that pushes $\rho(\phi)$ down. Compared to a population of only Romerians, the heterogeneous mix delivers a larger contraction region and a smaller multiplicity region. Pessimists, paradoxically, stabilize the equilibrium structure of the economy.

¹⁴The construction of monotone pure best responses uses the single-crossing approach of Athey (2001), which delivers a non-empty, compact, convex-valued, closed-graph best-response correspondence to which Kakutani's theorem can be applied.

Table 5. Equilibrium Regions and Their Economic Content

Region	Condition	$\rho(\phi)$	Equilibrium	Where it applies
Malthusian-dominated	$\phi < \phi^*$	< 0	Unique unconditionally	Trap; pre-industrial era
Contraction	$\phi^* \leq \phi \leq \phi^{**}$	$[0, 1)$	Unique via contraction	Most calibrations
Multiplicity	$\phi > \phi^{**}$	> 1	Multiple, selection by GG	Strong scale effects

Notes: Multiplicity region is non-empty iff $\theta_R > \theta_R^{crit}$, with $\theta_R^{crit} \approx 1.23$ at the baseline calibration $(\beta, \gamma) = (0.30, 0.325)$ and $n_R^* = 1.2$. GG denotes the global-games selection of Theorem 2.

5 Long-Run Dynamics and the Role of Wrong Beliefs

I now analyze the long-run dynamics of beliefs and population. Wrong beliefs generate qualitatively different pathologies depending on which worldview is correct. In Romerian worlds, Malthusian minorities cause prolonged stagnation through an informational poverty trap. In Malthusian worlds, Romerian minorities generate non-monotone dynamics, including the boom-bust cycles familiar from pre-industrial demography. The asymmetry arises from the interaction between demographic selection (which always favors optimists) and Bayesian learning (which favors correct beliefs, but only when signals are informative).

5.1 Two Forces: Demographic Selection and Learning

The Romerian share ϕ_t evolves according to (13), driven by two forces.

Demographic selection. Since $n_R > n_M$, $\phi_t^{birth} > \phi_t$ for any interior ϕ_t . Demographic selection pushes the cultural-transmission part of ϕ_{t+1} upward regardless of which belief is correct.

Bayesian learning. The direction of π_{t+1} depends on the true θ^* . The *strength* of learning depends on population through the signal informativeness $\mathcal{I}_t$ in (12). When $N_t \approx \bar{N}$, $\mathcal{I}_t \approx 0$ and learning is weak regardless of the truth.

In Romerian worlds, both forces align toward correct beliefs; the only question is how quickly correct beliefs spread, and the answer can be “arbitrarily slowly” when the economy starts in the trap. In Malthusian worlds, the two forces conflict, and the resulting tension generates the non-monotone dynamics that the rest of this section makes precise.

5.2 Dynamics in a Romerian World: Convergence and the Trap

Theorem 3 (Convergence in Romerian Worlds). *Suppose $\theta^* > 0$ and that the population trajectory remains within the region where (3) is well-defined. Starting from any interior*

initial condition $(\phi_0, N_0, A_0, \pi_0)$: (i) $\pi_t \rightarrow 1$ almost surely; (ii) $\phi_t \rightarrow 1$ almost surely; (iii) aggregate fertility converges to n_R^* ; (iv) $\sup_t \ln(N_t/\bar{N}) = +\infty$ almost surely.

The proof is in Appendix A. Both demographic selection and learning favor Romerians when the world is truly Romerian.¹⁵

Theorem 3 guarantees almost-sure convergence but is silent on the rate. The next analysis shows that convergence can be arbitrarily slow.

The Uninformative Region.

Definition 2 (Uninformative Region). *For $\delta > 0$, the δ -uninformative region is $\{N_t : |\ln(N_t/\bar{N})| < \delta\}$. In this region, the per-period information bound is $\mathcal{I}_t < (\theta_R - \theta_M)^2 \delta^2 / (2\sigma_\eta^2) \equiv \bar{\mathcal{I}}(\delta)$.*

When δ is small, signals reveal almost nothing about the true θ^* .

Definition 3 (Trap Region). *The informational poverty trap region $\mathcal{T}(\delta, \zeta)$ is the set of states (ϕ, N, π) with: (i) $|\ln(N/\bar{N})| < \delta$; (ii) $\bar{n}(\phi) \in [1, 1 + \zeta]$ for small $\zeta > 0$; and (iii) $\pi \leq \phi$, so posterior pessimism aligns with prevalent Malthusianism.*

The third condition ensures the trap is closed under the joint dynamics of (ϕ_t, N_t, π_t) : once $\pi_t \leq \phi_t$, this inequality propagates inductively (see proof of Proposition 2).

Assumption 4 (Fertility Straddles Replacement). *Malthusian fertility is below replacement and Romerian fertility above: $n_M < 1 < n_R$.*

Under Assumption 4, there is a critical Romerian share $\phi^{trap} = (1 - n_M)/(n_R - n_M)$ at which $\bar{n} = 1$. For ϕ slightly above ϕ^{trap} , aggregate fertility slightly exceeds 1 and population grows slowly. ϕ^{trap} lies within the Malthusian-dominated or contraction regions for standard calibrations.

Proposition 2 (Slow Learning in the Trap). *Consider a Romerian world ($\theta^* = \theta_R > 0$) under Assumption 4. Suppose $(\phi_0, N_0, \pi_0) \in \mathcal{T}(\delta, \zeta)$ for small $\delta, \zeta > 0$. For target $\bar{\phi} > \phi_0$, define the escape time $T^{esc} = \inf\{t : \phi_t > \bar{\phi}\}$. Then, as long as the trajectory remains in $\mathcal{T}(\delta, \zeta)$:*

$$(i) \quad |\mathbb{E}[\pi_{t+1} - \pi_t \mid \mathcal{F}_t]| \leq C_1 \bar{\mathcal{I}}(\delta) \text{ with } C_1 = 1/4;$$

¹⁵The result that $\sup_t \ln(N_t/\bar{N}) = +\infty$ presumes the absence of endogenous bounds on population. The framework can be extended to deliver bounded population through a hard carrying capacity or endogenous congestion. Under either extension, $\phi_t \rightarrow 1$ almost surely, but population converges to a stochastic neighborhood of the carrying capacity. The informational poverty trap is unaffected because it operates at populations close to $\bar{N}$.

- (ii) $\phi_{t+1}^{birth} - \phi_t \leq C_2(n_R - n_M)(1 - \phi_t)$ with $C_2 = 1/n_M$;
- (iii) $\mathbb{E}[T^{esc}] \geq (\bar{\phi} - \phi_0)/D$, where $D = \lambda C_2(n_R - n_M) + (1 - \lambda)C_1\bar{\mathcal{I}}(\delta)$.

The proof is in Appendix A. As $\delta \rightarrow 0$, the learning contribution vanishes and escape relies on demographic selection, which is itself slow when $n_R - n_M$ or λ is small. Corollary 3 shows that the expected escape time diverges as $\Omega(1/\delta^2)$ when learning is the binding channel and can exceed any finite horizon.

Proposition 8 characterizes the boundary of the trap. The critical population beyond which learning dominates demographic selection is

$$N^{crit}(\phi)/\bar{N} = \exp\left(\sqrt{\frac{2\sigma_\eta^2 \cdot \lambda C_2(n_R - n_M)(1 - \phi)}{(1 - \lambda)C_1(\theta_R - \theta_M)^2}}\right).$$

High noise σ_η^2 raises N^{crit} , making the trap larger; strong cultural transmission λ also raises N^{crit} ; and large belief dispersion $(\theta_R - \theta_M)^2$ lowers N^{crit} by making signals more diagnostic.

In a Romerian world, then, Malthusian minorities cause prolonged stagnation through a self-reinforcing mechanism: pessimism suppresses fertility, low fertility keeps population near $\bar{N}$, and uninformative signals fail to correct pessimism. Escape eventually occurs but can take arbitrarily long.¹⁶

5.3 Dynamics in a Malthusian World: Opposing Forces and Non-Monotone Outcomes

The dynamics differ qualitatively when $\theta^* < 0$. Whereas in a Romerian world demographic selection and learning push ϕ_t in the same direction, in a Malthusian world they pull in opposite directions.

Proposition 3 (Opposing Forces in Malthusian Worlds). *Suppose $\theta^* = \theta_M < 0$. At any interior state with $\phi_t \in (0, 1)$ and $\pi_t \in (0, 1)$:*

- (i) Demographic-selection drift on births is strictly positive and bounded. *The contribution of differential fertility to the share of Romerian-parented children satisfies*

$$0 < \phi_t^{birth} - \phi_t = \frac{\phi_t(1 - \phi_t)(n_R - n_M)}{\bar{n}_t} \leq \frac{n_R - n_M}{4n_M},$$

independently of $\ln(N_t/\bar{N})$.

¹⁶This mechanism differs sharply from settings where delay is informative: in some dynamic-learning models (e.g., Gans, 2026), silence carries information and high variance only postpones rather than prevents action. Here, the analog of silence—population at the reference $\bar{N}$ —is by construction *uninformative* about θ^* , so trap durations can be unboundedly large rather than merely lengthy.

(ii) Bayesian learning drift on the posterior is strictly negative under truth and grows quadratically in population deviation. *The expected change in the posterior under the true law of motion is*

$$\mathbb{E}^{\theta_M}[\pi_{t+1} - \pi_t \mid \mathcal{F}_t] = -\pi_t(1 - \pi_t)\mathcal{I}_t + o(\mathcal{I}_t), \quad \mathcal{I}_t = \frac{(\theta_R - \theta_M)^2 [\ln(N_t/\bar{N})]^2}{2\sigma_\eta^2}.$$

(iii) The expected drift in ϕ_t admits a sign change as a function of population deviation. Using (13), the joint drift decomposes as

$$\mathbb{E}^{\theta_M}[\phi_{t+1} - \phi_t \mid \mathcal{F}_t] = \underbrace{\lambda(\phi_t^{birth} - \phi_t)}_{\geq 0, \text{ selection}} + \underbrace{(1 - \lambda)(\pi_t - \phi_t)}_{\text{cultural regression toward } \pi_t} + \underbrace{(1 - \lambda)\mathbb{E}^{\theta_M}[\pi_{t+1} - \pi_t \mid \mathcal{F}_t]}_{\leq 0, \text{ learning}}.$$

At $\ln(N_t/\bar{N}) = 0$ the learning term vanishes; for $|\ln(N_t/\bar{N})|$ sufficiently large, the learning term dominates and the drift is strictly negative whenever $\pi_t \in (0, 1)$. The deterministic skeleton admits no monotone law of motion in ϕ_t .

The proof, in Appendix A, is a direct Bayes log-odds calculation combined with a bounded-selection-drift comparison. The threshold takes the form

$$\bar{L}(\phi_t, \pi_t, \lambda)^2 = \frac{2\sigma_\eta^2 \cdot [\lambda(\phi_t^{birth} - \phi_t) + (1 - \lambda) \max\{\pi_t - \phi_t, 0\}]}{(1 - \lambda) \pi_t(1 - \pi_t) (\theta_R - \theta_M)^2},$$

which is decreasing in the Romerian–Malthusian fertility gap and increasing in signal noise.

Economic content. Proposition 3 captures the core economic asymmetry. In a Malthusian world, demographic selection always pushes upward because optimists outbreed pessimists; learning under truth always pushes downward because outcomes confirm the Malthusian model; and the magnitude of the learning channel scales quadratically with how far population has deviated from the reference. The drift is therefore non-negative near $\bar{N}$ and strictly negative once population has grown enough that the learning channel overtakes selection. Any candidate interior steady state must lie at a population strictly above the reference, where $\bar{g} + \theta_M \ln(N^{ss}/\bar{N}) = 0$ gives $\ln(N^{ss}/\bar{N}) = -\bar{g}/\theta_M > 0$.

Cycle dynamics. The state-dependent reversal naturally produces the cyclical pattern familiar from economic-historical accounts of the pre-industrial era, with three identifiable phases. *Expansion phase:* population is moderate, learning is weak, demographic selection dominates, and ϕ_t rises. *Crisis phase:* as ϕ_t and population grow, signals become informative; productivity slows because $\theta^* < 0$; the Malthusian model fits much better; π_{t+1} falls

sharply, overwhelming demographic selection; and ϕ_{t+1} falls. *Recovery phase*: fertility falls, population stabilizes or contracts, signals weaken, demographic selection regains the upper hand, and the pattern can repeat. Whether the system traces out a deterministic limit cycle, damped oscillations, or stochastic recurrence depends on curvature of the nonlinear dynamics and shock size.

Comparative statics for the strength of cycles. When learning is very weak, $\bar{L}$ is large and the drift remains positive over a wide population range; demographic selection dominates and ϕ_t rises persistently. When learning is very strong, $\bar{L}$ is small and the drift turns negative for modest population deviations; Malthusian beliefs spread quickly, and the system settles into a Malthusian-dominated regime. The most pronounced non-monotone dynamics arise for intermediate learning responsiveness, where neither force dominates globally.

5.4 The Asymmetry: Why Wrong Beliefs Generate Different Pathologies

Table 6 summarizes the asymmetry.

Table 6. Pathologies from Wrong Beliefs

	Romerian World ($\theta^* > 0$)	Malthusian World ($\theta^* < 0$)
Wrong belief	Malthusian ($\theta_M < 0$)	Romerian ($\theta_R > 0$)
Demographic selection	Favors Romerians (correct)	Favors Romerians (incorrect)
Learning	Favors Romerians (correct)	Favors Malthusians (correct)
Forces align?	Yes	No
Pathology	Informational poverty trap	Non-monotone dynamics, including cycles
Long-run outcome	Convergence to truth	State-dependent reversal

In Romerian worlds, both forces align toward correct beliefs. Incorrect (Malthusian) beliefs are self-correcting, but the correction can be very slow if the economy starts in the trap. In Malthusian worlds, the forces conflict: demographic selection favors optimists who happen to be wrong, while learning favors the correct pessimistic model.

This asymmetry has a profound implication. Human societies are more vulnerable to excessive optimism in terms of stability, but more vulnerable to excessive pessimism in terms of stagnation. Optimists reproduce more regardless of whether their optimism is warranted, which destabilizes Malthusian worlds. Pessimism keeps population near the reference and thereby suppresses learning, which traps Romerian worlds.

Remark 1 (The Core Insight). In a Romerian world, Malthusians cause stagnation; in a Malthusian world, Romerians generate non-monotone dynamics, including cycles. *Wrong beliefs are always costly, but the nature of the cost depends on the direction of the error.*

6 Demographic History: Two Interpretations

The model offers a unified framework for interpreting broad patterns in demographic history. It admits two fundamentally different interpretations of the pre-industrial era, with distinct implications for understanding the modern world.

6.1 Interpretation 1: A Genuinely Malthusian Past

The conventional interpretation, consistent with economic historians from Malthus himself to [Clark \(2007\)](#), is that the pre-industrial world was genuinely Malthusian: scale effects were negative, and larger populations did reduce living standards.

Under this interpretation, the cycles documented by [Voigtländer and Voth \(2013\)](#) and others reflect exactly the kind of non-monotone dynamics established in Proposition 3. During periods of relative prosperity, Romerian-like optimism spread through demographic selection; population grew; eventually, population deviation from the reference became large enough that the learning channel engaged; living standards fell, Malthusian beliefs were reinforced, and the drift in ϕ_t reversed.

The Black Death of 1347–1351 provides a particularly clear example. Europe’s population had grown substantially, straining resources. The plague killed roughly one-third of the population. Real wages rose sharply in the aftermath and remained elevated for over a century. The model interprets this as a crisis phase that both directly reduced population and reinforced Malthusian beliefs.

Under this interpretation, the Industrial Revolution marked a genuine regime change from $\theta^* < 0$ to $\theta^* > 0$. Sustained technological progress, enabled by the accumulation of scientific knowledge and supporting institutions, broke the Malthusian link.

6.2 Interpretation 2: An Always-Romerian World Trapped by Information Poverty

The informational poverty trap opens a radically different interpretation: the world may have *always* been Romerian, but humanity was trapped for millennia by the self-reinforcing dynamics of pessimism.

Under this interpretation, positive scale effects existed throughout history. But populations were too close to the relevant reference, and productivity signals too noisy, for anyone to reliably learn this fact. Malthusian beliefs, once established, suppressed fertility, keeping population near $\bar{N}$. The trap persisted not because Malthusianism was correct, but because the evidence needed to refute it could only be generated by escaping it.

The observed pre-industrial cycles, under this interpretation, were not the inevitable boom-bust dynamics of a Malthusian world but rather failed escape attempts from the trap. Occasionally, some region or era experienced conditions that pushed population higher: better weather, new crops, reduced warfare. Beliefs began to update toward optimism. But the gains were fragile. Any negative shock—plague, famine, invasion—was interpreted as confirmation of Malthusian fears, beliefs snapped back, and the economy returned to the trap.

The Industrial Revolution, under this interpretation, was an *informational escape*. Some combination of factors—the Enlightenment, improving institutions, a sequence of positive productivity shocks—pushed European populations above the critical threshold N^{crit} . Once in the informative region, learning accelerated; evidence of positive scale effects accumulated; beliefs shifted decisively toward Romerianism; fertility rose further; population exploded.

6.3 Distinguishing the Interpretations

Both interpretations are consistent with the qualitative features of demographic history. Several features favor the genuine-regime-change interpretation: the timing of the Industrial Revolution coincides with specific technological innovations; the scale of progress after 1800 suggests a qualitative break; environmental constraints appear to have genuinely bound pre-industrial economies. Several features favor the informational-escape interpretation: there is evidence of positive scale effects even in pre-industrial settings; the speed of the transition is more consistent with belief updating than gradual technological change; the modern fertility decline is puzzling under a strict Romerian view but explicable if beliefs can drift back toward Malthusianism.

I do not take a definitive stand. The contribution of the model is to show that *both are theoretically coherent* and that they have *different implications for policy*.

6.4 The Modern Fertility Decline

Fertility rates in developed countries have fallen dramatically since the 1960s, reaching levels well below replacement in most of Europe and East Asia. The model suggests several interpretations.

Revival of Malthusian beliefs. The environmental movement can be understood as a revival of Malthusian thinking. The empirical evidence in Section 3 documents that such beliefs are prevalent and associated with substantially lower fertility. Under the informational-escape interpretation, the modern fertility decline represents a partial return toward the

trap.

Actual shift to Malthusian technology. Under the regime-change interpretation, the true θ^* may have recently shifted from positive to negative. Automation and AI may have reduced the marginal contribution of human labor; environmental constraints may impose genuine costs on population growth.

Quantity-quality substitution. The model abstracts from human capital, but in reality parents face a quantity-quality tradeoff. Rising returns to education may have induced parents to invest more in each child’s human capital. This is conceptually distinct from Malthusianism but generates similar fertility outcomes.

Under the dynastic welfare criterion, what matters is whether households’ beliefs match the truth. Under the total utilitarian criterion, the planner additionally values population for its own sake.

6.5 Contemporary Divergence and Future Dynamics

Fertility today varies enormously across countries and groups. Within countries, fertility varies systematically with religion, education, and political orientation. The model predicts that these differences are partially driven by belief heterogeneity and are self-reinforcing through demographic selection. Standard demographic projections typically assume convergence to low fertility as countries develop, but the model suggests convergence may be slower or incomplete if belief heterogeneity persists. Whether this is efficient depends on the true value of θ^* .

7 Policy: Welfare Criteria and the Value of Public Information

I now analyze the welfare effects of public announcements about the scale-effect parameter θ . The analysis combines an information-provision channel that operates throughout the share space and a selection-shifting channel that operates inside the multiplicity region. A third subsection introduces a private-versus-public information distinction, with a closed-form retained-fraction formula governing how much of a public announcement is dissipated through equilibrium feedback.

7.1 Welfare Criteria with Endogenous Population

Evaluating welfare when fertility is a choice variable raises fundamental conceptual issues: the set of people who exist depends on policy.

Dynastic Welfare. Under the dynastic criterion, the planner respects households' subjective preferences:

$$W^{dyn} = \phi_0 V_R(A_0, N_0, \phi_0, \pi_0) + (1 - \phi_0) V_M(A_0, N_0, \phi_0, \pi_0). \quad (23)$$

This is closely related to the P-efficiency concept of [Goloso et al. \(2007\)](#).

Total Utilitarian Welfare.

$$W^{tot} = \mathbb{E}^* \left[\sum_{t=0}^{\infty} \beta^t N_t \cdot \ln c_t \right]. \quad (24)$$

For (24) to be well-defined, $\beta \mathbb{E}^* n_R^* < 1$ at the relevant equilibrium fertility. Under the time-cost specification, $\ln c_t = \ln A_t + \ln(1 - \psi(n_t))$ decomposes additively.

Average utilitarian benchmark.

$$W^{avg} = \mathbb{E}^* \left[\sum_{t=0}^{\infty} \beta^t \ln c_t \right]. \quad (25)$$

W^{avg} internalizes the per-capita scale-effect channel but omits the *population-multiplier* term N_t .

The dynastic and total utilitarian criteria span a substantively distinct normative spectrum. The dynastic planner is fundamentally an information provider; the total utilitarian planner is an externality corrector. The average-utilitarian benchmark sits between them.

7.2 The First-Best Allocation and Welfare Decomposition

A planner who knows θ^* and can directly choose a symmetric fertility path under the total utilitarian criterion solves

$$\max_{\{n_t\}} \mathbb{E}^* \left[\sum_{t=0}^{\infty} \beta^t N_t \cdot (\ln A_t + \ln(1 - \psi(n_t))) \right] \quad (26)$$

subject to $N_{t+1} = n_t N_t$ and (3). Conjecturing $V^P(A, N) = N \cdot w^P(A, N)$ with $w^P(A, N) \approx v^P + \xi_A^P \ln A + \xi_N^P \ln(N/\bar{N})$, the planner's FOC for fertility is

$$\frac{\psi'(n_t^{FB})}{1 - \psi(n_t^{FB})} = \beta [w^P(A', N') + \xi_N^P]. \quad (27)$$

The household's common-correct-belief FOC is structurally identical except for the absence of ξ_N^P . The wedge between planner and household optimality is exactly $\beta\xi_N^P$, with the sign of θ^* .

Proposition 4 (First-Best vs. Decentralized Fertility). *In a Romerian world ($\theta^* > 0$): (i) $n^{FB} > n^{CB}$, where n^{CB} is the common-correct-belief decentralized fertility; (ii) $n^{FB} > n^{CB} > n_M^*$; (iii) the welfare gap decomposes into externality, mean-bias, and dispersion components. In a Malthusian world ($\theta^* < 0$), $n^{FB} < n^{CB}$.*

The proof is in Appendix B.

Proposition 5 (Welfare Decomposition). *Under the total utilitarian criterion,*

$$W^{FB}(\theta^*) - W^{HB} = \underbrace{\Delta_1(\theta^*)}_{\text{externality}} + \underbrace{\Delta_2(\theta^*, \bar{\theta})}_{\text{mean bias}} + \underbrace{\Delta_3(\text{Var}(\theta_i))}_{\text{dispersion}}, \quad (28)$$

where $\Delta_1 \geq 0$ with equality iff $\theta^* = 0$; $\Delta_2 \geq 0$ is increasing in $(\bar{\theta} - \theta^*)^2$; and $\Delta_3 \approx (\omega\kappa^2/2) \text{Var}(\theta_i)$.¹⁷

The proof is in Appendix B. Public announcements can reduce Δ_2 and Δ_3 but cannot directly address Δ_1 without additional policy instruments. The decomposition delivers three messages. First, even with $\bar{\theta} = \theta^*$ so $\Delta_2 = 0$, the externality Δ_1 persists. Second, even with no dispersion and correct mean beliefs, the externality remains. Third, dispersion generates welfare loss even when $\bar{\theta} = \theta^*$.

7.3 The Planner's Announcement Problem: Two Channels

Consider a planner who observes a noisy estimate of the true parameter,

$$\hat{\theta} = \theta^* + \nu, \quad \nu \sim \mathcal{N}(0, \sigma_\nu^2), \quad (29)$$

and publicly releases a signal $z = \mu + \tilde{\nu}$, where $\tilde{\nu} \sim \mathcal{N}(0, \tau_A^{-1})$. The announcement operates through two channels.

(1) *Information-provision channel.* Households use z in their Bayes update over $\{\theta_R, \theta_M\}$, shifting the posterior π_t^{post} toward the announced value.

(2) *Selection-shifting channel.* In the multiplicity region, the announcement also serves as a public signal about which equilibrium will be played. Through Theorem 2, an announcement can move the cutoff ϕ^{rd} and switch the selection between low-fertility and high-fertility stable equilibria.

¹⁷The Taylor expansion underlying Δ_3 is taken around the social-optimum fertility n^{SP} , so the externality Δ_1 appears explicitly rather than being absorbed.

The two channels reinforce each other in Romerian worlds with strong scale effects.

Assumption 5 (Announcement Credibility). *Households treat the announcement as a Gaussian signal with precision $\tau_A > 0$. When $\tau_A \rightarrow \infty$, households fully adopt the announced value; when $\tau_A \rightarrow 0$, households ignore the announcement.*

7.4 Optimal Announcement under Different Criteria

Theorem 4 (Optimal Announcement: Dynastic Welfare). *Under W^{dyn} , the optimal announcement is truthful in expectation: $\mu^{dyn} = \hat{\theta}$, and the planner prefers maximal informativeness.*

Under the dynastic criterion, the planner is fundamentally an information provider. There is no externality the planner cares about beyond what households themselves would, so distorting the announcement strictly reduces welfare.

Theorem 5 (Optimal Announcement: Total Utilitarian Welfare). *Consider $\theta^* = \theta_R > 0$. Under W^{tot} , the optimal announcement satisfies*

$$\mu^{tot} = \hat{\theta} + \chi^{tot}, \quad (30)$$

where $\chi^{tot} > 0$ is an “optimism shade” whenever the externality $\Delta_1 > 0$. The shade has the following comparative statics: (i) χ^{tot} is increasing in $|\theta^*|$; (ii) χ^{tot} is increasing in $|\ln(N_0/\bar{N})|$; (iii) χ^{tot} is decreasing in $\text{Var}(\theta_{i,0})$; (iv) χ^{tot} is increasing in announcement credibility τ_A . In a Malthusian world, the sign of χ^{tot} reverses.

The proof is in Appendix B. The intuition is that the planner equates the marginal utilitarian benefit of shifting average fertility (carrying an externality wedge of size $\beta\xi_N^P$) to the marginal cost of distorting individual fertility decisions away from their belief-conditional optimum.

Corollary 1 (Optimal Announcement: Average Utilitarian Welfare). *Under W^{avg} , the optimal announcement is $\mu^{avg} = \hat{\theta} + \chi^{avg}$, where $0 < \chi^{avg} < \chi^{tot}$ in a Romerian world (with the inequalities reversing in a Malthusian world). The wedge driving χ^{avg} is the per-capita scale-effect spillover $\beta\xi_{N,avg}^P$, which is strictly smaller in absolute value than $\beta\xi_N^P$ because it omits the population-multiplier term N_t .*

Corollary 2 (Ordering of Announcements). *In a Romerian world, $\mu^{dyn} = \hat{\theta} < \mu^{avg} < \mu^{tot}$, with the inequalities reversing in a Malthusian world.*

7.5 Private versus Public Information: A Retained-Fraction Result

The announcement z studied above is public: it shifts every household's prior by the same amount. In practice, planners can also operate through private channels—parental belief transmission, household-specific advisory contacts, targeted demographic guidance. The two channels are not equivalent.

Consider two stylized policy instruments, both equally informative about θ^* . The first, *private validation* of magnitude q , shifts the prior of a single dynasty from its current value to that value plus q , leaving all other households' priors unchanged. The second, *public validation* of magnitude z , shifts every household's prior by z . A planner who controls one or the other faces different welfare consequences, even when the magnitudes are equal.

Proposition 6 (Retained Fraction of Public Information). *In a Romerian world, suppose all households share a common prior shifted by a public-validation increment z , and let $\bar{n}^{post}(z)$ denote the resulting equilibrium aggregate fertility. To first order around the balanced-growth reference, the dynastic welfare gain from a unit increment of public validation is*

$$\left. \frac{dV_R^{dyn}}{dz} \right|_{public} = \underbrace{\frac{2a}{2a + \Delta_{wedge}}}_{\text{retained fraction}} \cdot \left. \frac{dV_R^{dyn}}{dq} \right|_{private}, \quad (31)$$

where $a > 0$ parameterizes the half-width of the rival-prior support and $\Delta_{wedge} > 0$ is the magnitude of the strategic-complementarity wedge induced by the equilibrium response of other Romerians. The retained fraction lies strictly in $(0, 1)$ and is decreasing in Δ_{wedge} .

The proof is in Appendix B. A public announcement raises a Romerian household's own optimism but, because all other Romerians are also shifting upward, the marginal continuation value of a child is partly competed away. A unit of private validation is not subject to this dissipation.

Comparative statics of the retained fraction. First, the retained fraction is decreasing in Δ_{wedge} , which is itself increasing in $|\theta^*|$. In environments with strong scale effects, public communication is comparatively less effective per unit of credibility. Second, the retained fraction is increasing in the half-width a of rival-prior support: belief dispersion is a buffer that protects from dissipation effects. Third, the dissipated fraction approaches one as the wedge becomes large relative to dispersion, so in extreme cases the public announcement contributes essentially nothing.

Policy implications. Per unit of credibility cost, private channels of belief transmission strictly dominate public channels. Examples include household-specific advising at points of major life transitions, parental information transmission through targeted parenting-resource interventions, and confidential demographic-policy consultations. Examples of public channels include broad media campaigns, open data dashboards, and prominent policy speeches.

Connection to the trap. Public announcements work best precisely when households are in the trap, because Δ_{wedge} is small there: the slope $\rho(\phi)$ is below one in the trap region. The retained fraction is therefore close to one. Public information policy is most effective at exactly the time it is most needed.

7.6 The Selection-Shifting Channel and the Planner’s Dilemma

When the economy is in or near the multiplicity region, the announcement operates additionally through the selection channel. The cutoff ϕ^{rd} is a function of expected fertility levels under each selection, both of which respond to households’ subjective θ^{post} . An optimism-shaded announcement therefore not only raises Romerian fertility on the high-fertility selection but also lowers the cutoff at which the high-fertility selection is played.

Proposition 7 (Selection-Shifting Effect). *In the multiplicity region, the optimal total utilitarian announcement satisfies $\mu^{tot,sel} = \mu^{tot} + \chi^{sel}$, where $\chi^{sel} > 0$ in a Romerian world. The magnitude of χ^{sel} is non-monotone in ϕ : it peaks when ϕ is near ϕ^{rd} and vanishes when ϕ is far from it on either side.*

The proof is in Appendix B. When the economy is firmly in the high-fertility basin, shifting the cutoff has no effect on the realized selection; when firmly in the low-fertility basin, the cutoff is far from binding; selection-shifting matters only at the margin.

A key insight emerges from the structure of the learning problem: the planner faces the same informational constraints as households. If the economy is in the uninformative region, the historical productivity signals that inform $\hat{\theta}$ are themselves uninformative. The planner’s signal precision is endogenously low precisely when the trap is most binding.

Remark 2 (The Planner’s Dilemma). *Public information is most valuable in the informational poverty trap, where coordinating beliefs could push population above N^{crit} and initiate escape. But this is precisely when the planner has the least information. The cases where intervention matters most are the cases where the planner is most uncertain about what to say.*

This creates a policy dilemma distinct from [Morris and Shin \(2002\)](#): there, the concern is that public information may be *too influential*; here, the planner may have *nothing useful to say* when saying something would matter most. In the informational poverty trap, even a moderately noisy announcement can improve welfare by breaking self-reinforcing pessimism. In a Malthusian world with non-monotone dynamics, announcements require higher precision to avoid amplifying volatility.¹⁸

Trap-region value. Appendix B, Proposition 9, shows that in the trap region in a Roman world, the expected value of truthful announcement under W^{tot} is $\mathcal{V}^{trap} = P^{esc} \cdot \Delta W^{esc} - P^{deepen} \cdot \Delta W^{deepen}$, with a precision threshold above which $\mathcal{V}^{trap} > 0$. Because ΔW^{esc} is large, even a moderately noisy signal can be worth announcing.

Non-monotone-region value. Proposition 10 establishes that in a Malthusian world, the value of announcement comes from reducing belief dispersion (and thereby dampening oscillations). This benefit is partially offset by adjustment costs from disrupting the existing dynamics.

The stabilizing role of Malthusians. In response to a public announcement, Romanians coordinate (amplifying the signal), but Malthusians anti-coordinate (dampening it). This heterogeneity reduces the risk of over-coordination on noise. The precision threshold for beneficial announcement is therefore *lower* in populations with substantial Malthusian minorities.

7.7 Comparison of Decentralized and Planner Allocations

Table 7 summarizes the contrast across criteria.

Under all three criteria, the decentralized equilibrium is inefficient due to the uninternalized scale-effect externality. The dynastic planner can address belief distortions but not the externality. The total utilitarian planner shades further, reflecting an additional motive to value population. Achieving full efficiency would require additional instruments (fertility subsidies in Romanian worlds, fertility taxes in Malthusian worlds).

7.8 Policy Implications

The analysis yields several implications for government communication policy.

¹⁸A further parallel deserves mention: anticipated ex-post pronatalist support during fertility crises creates a time-inconsistency problem analogous to bailout policies. Anticipating support raises ex-ante fertility above the level optimal without commitment, amplifying cycle dynamics in a Malthusian world. Credible long-run population policy requires explicit commitment limits.

Table 7. Decentralized vs. Planner Allocations by Welfare Criterion

	Dynastic	Average Util.	Total Util.
<i>Romerian World ($\theta^* > 0$)</i>			
Decentralized fertility		Too low (externality + wrong beliefs)	
Optimal announcement	Truthful	Mildly shaded toward θ_R	Strongly shaded toward θ_R
Selection-shifting role	None	Modest	Strongest
Residual inefficiency	Externality only	Reduced externality	Most reduced externality
<i>Malthusian World ($\theta^* < 0$)</i>			
Decentralized fertility		Too high (externality), mixed by beliefs	
Optimal announcement	Truthful	Mildly shaded toward θ_M	Strongly shaded toward θ_M
Selection-shifting role	None	Modest	Strongest
Residual inefficiency	Externality only	Reduced externality	Most reduced externality

In the trap, take calculated risks. When the economy appears trapped, even a moderately imprecise signal may be worth announcing. The retained-fraction result reinforces this: in the trap, the equilibrium-response wedge is small, so public information transmits efficiently into action.

Under dynastic welfare, be truthful. If the planner respects household preferences as given, the optimal policy is to provide accurate information.

Under utilitarian criteria, shade systematically. The total utilitarian planner shades announcements toward optimism in Romerian worlds and toward pessimism in Malthusian worlds, by an amount that scales with the externality, the deviation of population from the reference, and credibility.

Prefer private channels to public ones, especially when externalities are large. Per unit of credibility, private channels outperform public channels by the retained fraction $2a/(2a + \Delta_{\text{wedge}})$. In environments with strong scale effects, the planner should rely disproportionately on parental and household-level information transmission.

Near the multiplicity boundary, the announcement does double duty. When ϕ is near ϕ^{**} , the announcement's selection-shifting role becomes economically meaningful in addition to its information-provision role.

During Malthusian non-monotone dynamics, time announcements carefully. Announcing during an expansion—when population is high and signals are informative—can

accelerate learning and dampen the swings.

Build credibility for when it matters; recognize the limits of knowledge. Planners should invest in credibility during normal times so that announcements are effective when intervention matters most. The most important policy implication, however, is epistemic humility. When the economy is in the trap, the planner knows the least about what to say. Acknowledging this uncertainty, rather than projecting false confidence, may be the most honest policy.

8 Conclusion

Beliefs about the economic consequences of population are not background assumptions in fertility models. They are first-order determinants of demographic and macroeconomic dynamics. This paper has formalized that observation in an overlapping generations model with heterogeneous beliefs about scale effects and a time cost of children, with multiplicative log-normal productivity dynamics that admit a tractable log-linear approximation around a balanced-growth reference. The resulting feedback between beliefs, fertility, and the informativeness of data produces a sharp duality. Malthusian minorities trap Romerian economies in stagnation, while Romerian minorities generate non-monotone dynamics in Malthusian economies. Wrong beliefs are always costly, but the form of the cost is determined by the truth they fail to track.

The strategic structure of the fertility game partitions the share space into three economically distinct regions: a Malthusian-dominated region of unconditional uniqueness, a contraction region of conditional uniqueness covering most empirically relevant calibrations, and a Romerian-dominated multiplicity region whose existence depends on the strength of scale effects. A global-games refinement, applied to a reduction of the game in which the Malthusian best response has been substituted out, selects a risk-dominant equilibrium in the multiplicity region. The trap region sits inside the contraction region for any standard calibration, so the trap intuition is robust to the multiplicity question. But where multiplicity is present, the planner’s announcement plays a selection-shifting role in addition to its information-provision role.

The informational poverty trap is the paper’s most distinctive theoretical object. Unlike traps grounded in technological nonconvexities or credit frictions, it arises endogenously from the geometry of Bayesian learning: signals are uninformative precisely in the region where pessimism keeps the economy confined, so the very beliefs that cause stagnation also insulate themselves from correction. This reframes a long-standing question in economic

history. The pre-industrial Malthusian era can be read as a true Malthusian regime that ended with a technological break, or as a Romerian regime trapped for millennia by self-reinforcing pessimism that finally dissolved when population breached the informative region. The two readings carry sharply different implications for whether the modern fertility decline represents a benign demographic adjustment or a slow drift back into the trap.

Belief heterogeneity also forces a reckoning with how welfare itself is defined. Because the population over which welfare is evaluated is endogenous to policy, the dynastic, average-utilitarian, and total utilitarian criteria recommend systematically different communication strategies, with shading magnitudes that strictly increase as the planner's normative commitments expand. The dynastic planner is fundamentally an information provider; the total utilitarian planner additionally values population. The same logic produces an unconventional message about public information: the planner is most uncertain exactly when announcements would be most valuable. A retained-fraction result governs how much of a public announcement survives equilibrium feedback: per unit of credibility, private channels of belief transmission strictly dominate public channels, with the wedge growing in the magnitude of the scale-effect externality.

Several extensions are natural: a complete classification of the deterministic skeleton in Malthusian worlds; sharper causal identification of the belief-fertility link; a full common-source extension of the independence-of-priors structure; integration with the quantity-quality margin central to unified growth theory; a microfoundation for announcement credibility through repeated-game reputation; and a quantitative implementation that could in principle adjudicate between the two interpretations of pre-industrial history. Whether one reads human history as a Malthusian regime that ended or a Romerian regime that nearly never began, the lesson is the same: what societies believe about themselves shapes what they become, and what they become determines what they can learn.

References

- Agrawal, Ajay, Arnaldo Camuffo, Alfonso Gambardella, Joshua S Gans, Erin L Scott, and Scott Stern, “The Foundations of Bayesian Entrepreneurship,” in “Bayesian Entrepreneurship,” Cambridge, MA: MIT Press, 2025. Forthcoming.
- Angeletos, George-Marios and Jennifer La’O, “Sentiments,” *Econometrica*, 2013, *81* (2), 739–779.
- , Fabrice Collard, and Harris Dellas, “Quantifying Confidence,” *Econometrica*, 2018, *86* (5), 1689–1726.
- Ashraf, Quamrul and Oded Galor, “The “Out of Africa” Hypothesis, Human Genetic Diversity, and Comparative Economic Development,” *American Economic Review*, 2013, *103* (1), 1–46.
- Athey, Susan, “Single Crossing Properties and the Existence of Pure Strategy Equilibria in Games of Incomplete Information,” *Econometrica*, 2001, *69* (4), 861–889.
- Azariadis, Costas and Allan Drazen, “Poverty Traps,” *Quarterly Journal of Economics*, 1990, *105* (2), 501–526.
- Bakkensen, Laura A and Lint Barrage, “Going Underwater? Flood Risk Belief Heterogeneity and Coastal Home Price Dynamics,” *Review of Financial Studies*, 2022, *35* (8), 3666–3709.
- Bakkensen, Laura, Toàn Phan, and Tsz-Nga Wong, “Leveraging the Disagreement on Climate Change: Theory and Evidence,” *Journal of Political Economy*, 2025, *133* (10), 3132–3166.
- Baldauf, Markus, Lorenzo Garlappi, and Constantine Yannelis, “Does Climate Change Affect Real Estate Prices? Only If You Believe in It,” *Review of Financial Studies*, 2020, *33* (3), 1256–1295.
- Banerjee, Abhijit V and Andrew F Newman, “Occupational Choice and the Process of Development,” *Journal of Political Economy*, 1993, *101* (2), 274–298.
- Barro, Robert J and Gary S Becker, “Fertility Choice in a Model of Economic Growth,” *Econometrica*, 1989, *57* (2), 481–501.
- Becker, Gary S, “An Economic Analysis of Fertility,” in “Demographic and Economic Change in Developed Countries,” Columbia University Press, 1960, pp. 209–240.
- Bernstein, Asaf, Matthew T Gustafson, and Ryan Lewis, “Disaster on the Horizon: The Price Effect of Sea Level Rise,” *Journal of Financial Economics*, 2019, *134* (2), 253–272.
- Bisin, Alberto and Thierry Verdier, “The Economics of Cultural Transmission and the Dynamics of Preferences,” *Journal of Economic Theory*, 2001, *97* (2), 298–319.
- and —, “The Economics of Cultural Transmission and Socialization,” in “Handbook of Social Economics,” Vol. 1, Elsevier, 2011, pp. 339–416.

- Blackorby, Charles, Walter Bossert, and David Donaldson, *Population Issues in Social Choice Theory, Welfare Economics, and Ethics*, Cambridge University Press, 2005.
- Bloom, Nicholas, Charles I Jones, John Van Reenen, and Michael Webb, “Are Ideas Getting Harder to Find?,” *American Economic Review*, 2020, 110 (4), 1104–1144.
- Bordalo, Pedro, Nicola Gennaioli, and Andrei Shleifer, “Diagnostic Expectations and Credit Cycles,” *Journal of Finance*, 2018, 73 (1), 199–227.
- Carlsson, Hans and Eric Van Damme, “Global Games and Equilibrium Selection,” *Econometrica*, 1993, 61 (5), 989–1018.
- Cavalli-Sforza, Luigi Luca and Marcus W Feldman, *Cultural Transmission and Evolution: A Quantitative Approach*, Princeton University Press, 1981.
- Clark, Gregory, *A Farewell to Alms: A Brief Economic History of the World*, Princeton University Press, 2007.
- Daruich, Diego, “The Macroeconomic Consequences of Early Childhood Development Policies,” Working Paper 2018-029, Federal Reserve Bank of St. Louis 2018.
- de la Croix, David and Matthias Doepke, “Inequality and Growth: Why Differential Fertility Matters,” *American Economic Review*, 2003, 93 (4), 1091–1113.
- Doepke, Matthias and Fabrizio Zilibotti, *Love, Money, and Parenting: How Economics Explains the Way We Raise Our Kids*, Princeton University Press, 2019.
- Felin, Teppo and Todd R Zenger, “Entrepreneurs as Theorists: On the Origins of Collective Beliefs and Novel Strategies,” *Strategic Entrepreneurship Journal*, 2009, 3 (2), 127–146.
- and —, “The Theory-Based View: Economic Actors as Theorists,” *Strategy Science*, 2017, 2 (4), 258–271.
- Frankel, David M, Stephen Morris, and Ady Pauzner, “Equilibrium selection in global games with strategic complementarities,” *Journal of Economic Theory*, 2003, 108 (1), 1–44.
- Galor, Oded, *Unified Growth Theory*, Princeton University Press, 2011.
- and David N Weil, “Population, Technology, and Growth: From Malthusian Stagnation to the Demographic Transition and Beyond,” *American Economic Review*, 2000, 90 (4), 806–828.
- Gans, Joshua S, “Belief Dispersion and Entrepreneurial Entry,” NBER Working Paper 35091, National Bureau of Economic Research 2026.
- Geanakoplos, John, “The Leverage Cycle,” *NBER Macroeconomics Annual*, 2010, 24 (1), 1–66.
- Golosov, Mikhail, Larry E Jones, and Michèle Tertilt, “Efficiency with Endogenous Population Growth,” *Econometrica*, 2007, 75 (4), 1039–1071.

- Howe, Peter D, Matto Mildenberger, Jennifer R Marlon, and Anthony Leiserowitz, “Geographic Variation in Opinions on Climate Change at State and Local Scales in the USA,” *Nature Climate Change*, 2015, 5 (6), 596–603.
- Jones, Charles I, “R&D-Based Models of Economic Growth,” *Journal of Political Economy*, 1995, 103 (4), 759–784.
- , “The End of Economic Growth? Unintended Consequences of a Declining Population,” *American Economic Review*, 2022, 112 (11), 3489–3527.
- Kalemli-Ozcan, Sebnem, “A Stochastic Model of Mortality, Fertility, and Human Capital Investment,” *Journal of Development Economics*, 2003, 70 (1), 103–118.
- Kremer, Michael, “Population Growth and Technological Change: One Million BC to 1990,” *Quarterly Journal of Economics*, 1993, 108 (3), 681–716.
- Lee, Ronald, “The Demographic Transition: Three Centuries of Fundamental Change,” *Journal of Economic Perspectives*, 2003, 17 (4), 167–190.
- Lee, Sang Yoon (Tim) and Ananth Seshadri, “On the Intergenerational Transmission of Economic Status,” *Journal of Political Economy*, 2019, 127 (2), 855–921.
- Malmendier, Ulrike and Stefan Nagel, “Depression Babies: Do Macroeconomic Experiences Affect Risk Taking?,” *Quarterly Journal of Economics*, 2011, 126 (1), 373–416.
- Malthus, Thomas Robert, *An Essay on the Principle of Population*, J. Johnson, London, 1798.
- Manuelli, Rodolfo E and Ananth Seshadri, “Explaining International Fertility Differences,” *Quarterly Journal of Economics*, 2009, 124 (2), 771–807.
- Morris, Stephen and Hyun Song Shin, “Unique Equilibrium in a Model of Self-Fulfilling Currency Attacks,” *American Economic Review*, 1998, 88 (3), 587–597.
- and —, “Social Value of Public Information,” *American Economic Review*, 2002, 92 (5), 1521–1534.
- and —, “Global Games: Theory and Applications,” in Mathias Dewatripont, Lars Peter Hansen, and Stephen J Turnovsky, eds., *Advances in Economics and Econometrics: Theory and Applications, Eighth World Congress*, Vol. 1, Cambridge University Press, 2003, pp. 56–114.
- Parfit, Derek, *Reasons and Persons*, Oxford University Press, 1984.
- Romer, Paul M, “Endogenous Technological Change,” *Journal of Political Economy*, 1990, 98 (5, Part 2), S71–S102.
- Scheinkman, José A and Wei Xiong, “Overconfidence and Speculative Bubbles,” *Journal of Political Economy*, 2003, 111 (6), 1183–1220.

- Simsek, Alp, “Speculation and Risk Sharing with New Financial Assets,” *Quarterly Journal of Economics*, 2013, 128 (3), 1365–1396.
- Voigtländer, Nico and Hans-Joachim Voth, “The Three Horsemen of Riches: Plague, War, and Urbanization in Early Modern Europe,” *Review of Economic Studies*, 2013, 80 (2), 774–811.

A Proofs and Auxiliary Theoretical Results

A note on the appendix proofs

Several appendix results below rely on the local log-linear approximation (15) of the household value function. Throughout, I distinguish between (a) parts of the argument that are exact under the multiplicative log-normal technology of (3), and (b) parts that hold to first order around a balanced-growth reference. The proof of Proposition 3 is a direct calculation that does not rely on a dynamical-systems reduction; the global-games selection of Theorem 2 adapts the supermodular-game machinery of Frankel et al. (2003) to the reduced Romerian subgame obtained by substituting out the Malthusian best response, with the existence step using the single-crossing argument of Athey (2001).

A.1 Proof of Lemma 1

Step 1: Bellman equation. Substituting $c_{j,t} = A_t(1 - \psi(n))$ and using $\ln c_{j,t} = \ln A_t + \ln(1 - \psi(n))$ into the dynastic objective (6) yields the Bellman equation (14):

$$V_j(A, N, \phi, \pi) = \max_{n \geq 0} \left\{ \ln A + \ln(1 - \psi(n)) + \beta n \mathbb{E}_j[V_j(A', N', \phi', \pi')] \right\}.$$

The choice variable n enters the period return only through $\ln(1 - \psi(n))$ and the discount-multiplier βn ; aggregate state variables enter parametrically.

Step 2: Local log-linear conjecture. I conjecture

$$V_j(A, N, \phi, \pi) = v_j(\phi, \pi) + \xi_{A,j} \ln A + \xi_{N,j} \ln(N/\bar{N}) + R_j(A, N, \phi, \pi), \quad (32)$$

where R_j is a remainder of order $O(\|(\ln A - \ln A^*, \ln(N/\bar{N}))\|^2)$ around a balanced-growth reference $(A^*, \bar{N})$. The remainder R_j is generated solely by the dependence of the optimal policy n_j^* on $(\ln A, \ln(N/\bar{N}))$ through the FOC; the technology side of the Bellman is exactly linear under (4), so it contributes no nonlinearity.

Step 3: Expected continuation under belief θ_j . Using (4),

$$\ln A' = \ln A + \bar{g} + \theta^* \ln(N/\bar{N}) + \eta_{t+1}.$$

Type- j households evaluate this expectation under their subjective belief θ_j and using $\mathbb{E}_j[\eta_{t+1}] = 0$:

$$\mathbb{E}_j[\ln A'] = \ln A + \bar{g} + \theta_j \ln(N/\bar{N}).$$

For population, with $N' = \bar{n}^e N$,

$$\mathbb{E}_j[\ln(N'/\bar{N})] = \ln(N/\bar{N}) + \ln \bar{n}^e.$$

Substituting into (32) and ignoring R_j at the reference,

$$\mathbb{E}_j[V'_j] = \mathbb{E}_j v_j(\phi', \pi') + \xi_{A,j} [\ln A + \bar{g} + \theta_j \ln(N/\bar{N})] + \xi_{N,j} [\ln(N/\bar{N}) + \ln \bar{n}^e] + O(\sigma_\eta^2), \quad (33)$$

where the $O(\sigma_\eta^2)$ term is a Jensen correction absorbed into v_j .

Step 4: Coefficient matching at the reference. Substituting (33) into the Bellman (14), evaluating at the optimal policy n_j^* (treated as locally constant in $(\ln A, \ln(N/\bar{N}))$):

$$v_j + \xi_{A,j} \ln A + \xi_{N,j} \ln(N/\bar{N}) = \ln A + \ln(1 - \psi(n_j^*)) + \beta n_j^* (\mathbb{E}_j v'_j + \xi_{A,j} [\ln A + \bar{g} + \theta_j \ln(N/\bar{N})] + \xi_{N,j} [\ln(N/\bar{N}) + \ln \bar{n}^e])$$

Matching coefficients on $\ln A$:

$$\xi_{A,j} = 1 + \beta n_j^* \xi_{A,j} \implies \xi_{A,j} = \frac{1}{1 - \beta n_j^*}.$$

Matching coefficients on $\ln(N/\bar{N})$:

$$\xi_{N,j} = \beta n_j^* (\xi_{A,j} \theta_j + \xi_{N,j}) \implies \xi_{N,j} (1 - \beta n_j^*) = \frac{\beta n_j^* \theta_j}{1 - \beta n_j^*},$$

yielding

$$\xi_{N,j} = \frac{\beta n_j^* \theta_j}{(1 - \beta n_j^*)^2}.$$

This gives (17). The constant term v_j absorbs $\bar{g} \xi_{A,j}$, $\xi_{N,j} \ln \bar{n}^e$, the Jensen term, and the residual $\beta n_j^* \mathbb{E}_j v'_j$. The condition $\beta n_j^* < 1$ ensures $\xi_{A,j}$ and $\xi_{N,j}$ are finite.

Step 5: Fertility FOC. The first-order condition with respect to n in (14):

$$\frac{\psi'(n_j^*)}{1 - \psi(n_j^*)} = \beta \mathbb{E}_j[V'_j].$$

Under the specification (5), $\psi'(n)/(1 - \psi(n)) = \gamma n$, and the FOC becomes

$$\gamma n_j^* = \beta \mathbb{E}_j[V'_j]. \quad (34)$$

Step 6: Existence and uniqueness of the fixed point. Substituting (33) into (34) and noting that $\xi_{A,j}$ and $\xi_{N,j}$ themselves depend on n_j^* , we obtain a fixed-point equation

$$\gamma n_j^* = \beta [\bar{v}_j(\phi, \pi) + \xi_{A,j}(n_j^*) (\ln A^* + \bar{g} + \theta_j \ln(N^*/\bar{N})) + \xi_{N,j}(n_j^*) (\ln(N^*/\bar{N}) + \ln \bar{n})]$$

at the reference $(A^*, N^*, \bar{n})$. Define the right-hand side as a function $F(n)$. $F(n)$ is continuous on $[0, 1/\beta)$, with $F(0)$ finite and $F(n) \rightarrow \infty$ as $n \uparrow 1/\beta$. The map $n \mapsto \gamma n - F(n)$ is therefore strictly increasing for n small enough and admits a unique root in $[0, \bar{n}^{\max}]$ for some $\bar{n}^{\max} < 1/\beta$. Equilibrium fertility n_j^* is uniquely determined by θ_j and the reference state.

Step 7: Comparative statics in θ_j . Differentiating (34) totally with respect to θ_j at fixed $(\bar{n}, A^*, N^*, \phi, \pi)$:

$$\gamma \frac{\partial n_j^*}{\partial \theta_j} = \beta \left[\frac{\partial \xi_{A,j}}{\partial n_j^*} \frac{\partial n_j^*}{\partial \theta_j} \cdot [\dots] + \xi_{A,j} \ln(N^*/\bar{N}) + \frac{\partial \xi_{N,j}}{\partial \theta_j} + \frac{\partial \xi_{N,j}}{\partial n_j^*} \frac{\partial n_j^*}{\partial \theta_j} \cdot [\dots] \right].$$

At a balanced-growth reference where $\ln(N^*/\bar{N}) = 0$, the explicit- θ_j term reduces to $\beta \partial \xi_{N,j} / \partial \theta_j = \beta \cdot \beta n_j^* / (1 - \beta n_j^*)^2 > 0$. Provided the implicit terms (in n_j^*) do not flip the sign—which is verified by direct calculation under $\beta n_j^* < 1$ —this gives $\partial n_j^* / \partial \theta_j > 0$. Writing $n_j^* = n^* h(\theta_j; \bar{n})$ with h increasing in θ_j follows by reorganization. In particular, $\theta_R > \theta_M$ implies $n_R^* > n_M^*$. $\square$

A.2 Proof of Proposition 1

For type j expecting aggregate fertility $\bar{n}^e$, the conjecture (32) and equation (33) give

$$\mathbb{E}_j[V_j'] = v_j^{\text{const}} + \xi_{A,j} \ln A + \xi_{A,j} \theta_j \ln(N/\bar{N}) + \xi_{N,j} \ln(N/\bar{N}) + \xi_{N,j} \ln \bar{n}^e.$$

Hence

$$\frac{\partial \mathbb{E}_j[V_j']}{\partial \bar{n}^e} = \frac{\xi_{N,j}}{\bar{n}^e}.$$

The FOC (34) reads $\gamma n_j^* = \beta \mathbb{E}_j[V_j']$. Differentiating both sides totally with respect to $\bar{n}^e$, holding all other state variables fixed:

$$\gamma \frac{dn_j^*}{d\bar{n}^e} = \beta \cdot \frac{\xi_{N,j}}{\bar{n}^e},$$

where the dependence of $\xi_{N,j}$ on n_j^* generates a second-order correction that vanishes at the reference. Solving and substituting the closed form for $\xi_{N,j}$:

$$\frac{dn_j^*}{d\bar{n}^e} = \frac{\beta \xi_{N,j}}{\gamma \bar{n}^e} = \frac{\beta^2 n_j^* \theta_j}{\gamma \bar{n}^e (1 - \beta n_j^*)^2}.$$

This is (18). The sign matches the sign of θ_j . □

A.3 Proof of Theorem 1

Equilibrium condition. Let

$$BR(\bar{n}^e) = \phi n_R^*(\bar{n}^e) + (1 - \phi) n_M^*(\bar{n}^e).$$

A rational-expectations equilibrium is a fixed point $\bar{n}^* = BR(\bar{n}^*)$. The local slope is

$$\rho(\phi) = BR'(\bar{n}^*) = \phi d_R + (1 - \phi) d_M,$$

with $d_R > 0$ and $d_M < 0$ by Proposition 1.

Part (i): Malthusian-dominated region ($\phi < \phi^*$). For $\phi < \phi^* = |d_M|/(d_R + |d_M|)$, $\rho(\phi) < 0$. BR is locally strictly decreasing in a neighborhood of any candidate fixed point. Define $\Phi(x) = BR(x) - x$, which is continuously differentiable with $\Phi'(x) = \rho(\phi) - 1 < -1 < 0$ at any fixed point. By the implicit function theorem, fixed points are isolated. Since both $n_R^*(\bar{n}^e)$ and $n_M^*(\bar{n}^e)$ are continuous and bounded above by $\bar{n}^{\max} < 1/\beta$, BR maps $[0, \bar{n}^{\max}]$ into itself; by the intermediate value theorem applied to Φ , at least one fixed point exists. Strict monotonicity of Φ implies the fixed point is unique. The argument requires only the sign of $\rho(\phi)$, not its magnitude.

Part (ii): Contraction region. For ϕ in this range, $0 \leq \rho(\phi) < 1$. The map $BR : [0, \bar{n}^{\max}] \rightarrow [0, \bar{n}^{\max}]$ is Lipschitz with constant $\rho(\phi) < 1$ on a neighborhood of any candidate fixed point. By Banach's fixed-point theorem, the fixed point is unique. When $d_R \leq 1$, $\phi^{**} \geq 1$, so the contraction region extends to all $\phi \in [\phi^*, 1]$.

Part (iii): Multiplicity region ($\phi > \phi^{}$; requires $d_R > 1$).** For $\phi > \phi^{**}$, $\rho(\phi) > 1$. Locally near any fixed point, $\Phi'(x) = \rho - 1 > 0$, so BR crosses the 45-degree line from below. Since BR is continuous on $[0, \bar{n}^{\max}]$ with $BR(0) \geq 0$ and $BR(\bar{n}^{\max}) \leq \bar{n}^{\max}$, Φ takes nonnegative values near 0 and nonpositive near $\bar{n}^{\max}$. By the intermediate value theorem

combined with the local slope condition, the number of transversal intersections is odd and at least three. Standard transversality arguments imply local slopes alternate between greater than one (unstable) and less than one (stable). The extreme intersections are stable.

Multiplicity condition. $\phi^{**} < 1$ iff $d_R > 1$. Substituting (18) for d_R gives

$$\frac{\beta^2 n_R^* \theta_R}{\gamma \bar{n} (1 - \beta n_R^*)^2} > 1 \iff \theta_R > \theta_R^{crit} \equiv \frac{\gamma \bar{n} (1 - \beta n_R^*)^2}{\beta^2 n_R^*}. \quad \square$$

A.4 Proof of Theorem 2

I adapt the supermodular global-games argument of Frankel et al. (2003) via a reduction that substitutes out the Malthusian best response. The existence step uses the single-crossing approach of Athey (2001).

Step 1: Reduction to a Romerian subgame. Let n_R^i denote individual Romerian fertility choices and let $\bar{n}_R = \int_R n_R^i di / \phi$ denote average Romerian fertility. The aggregate fertility identity is

$$\bar{n} = \phi \bar{n}_R + (1 - \phi) n_M^*(\bar{n}).$$

Because Malthusians have a strictly downward-sloping best response with slope $d_M < 0$, the map $\bar{n} \mapsto \phi \bar{n}_R + (1 - \phi) n_M^*(\bar{n})$ for fixed $\bar{n}_R$ has derivative $(1 - \phi) d_M < 0$, and so the equation has a unique solution $\bar{n} = T(\bar{n}_R)$. By the implicit function theorem,

$$T'(\bar{n}_R) = \frac{\phi}{1 - (1 - \phi) d_M} > 0.$$

The reduced game has Romerian players choosing n_R^i , with payoffs $\tilde{U}_R(n_R^i, \bar{n}_R) := U_R(n_R^i, T(\bar{n}_R))$.

Step 2: Supermodularity of the reduced game.

$$\frac{\partial^2 \tilde{U}_R}{\partial n_R^i \partial \bar{n}_R} = \frac{\partial^2 U_R}{\partial n_R^i \partial \bar{n}} \cdot T'(\bar{n}_R).$$

By Proposition 1, $\partial^2 U_R / \partial n_R^i \partial \bar{n} > 0$, and $T'(\bar{n}_R) > 0$. Therefore the reduced game is strictly supermodular in actions among Romerians.

Step 3: Monotone signal structure and existence of monotone pure best responses. Each Romerian observes $\phi_i = \phi + \zeta_i$ with $\zeta_i \sim \mathcal{N}(0, \sigma_\zeta^2)$ i.i.d. With a diffuse improper prior on ϕ , the conditional distribution of ϕ given ϕ_i is also Gaussian with mean

ϕ_i and variance σ_ζ^2 , satisfying the monotone likelihood ratio property. The reduced game now satisfies the conditions of Theorem 1 of [Athey \(2001\)](#): the payoff difference between any two actions is single-crossing in the type, the type space is ordered, and posteriors satisfy MLRP. Athey's theorem guarantees the existence of monotone pure-strategy equilibria.

Step 4: Dominance regions. Under the calibration of the multiplicity region, there exist values $\bar{\phi}_L < \phi^{**}$ and $\bar{\phi}_H > 1$ such that for $\phi < \bar{\phi}_L$ the low-fertility selection strictly dominates regardless of others' actions, and for $\phi > \bar{\phi}_H$ the high-fertility selection strictly dominates.

Step 5: Application of [Frankel et al. \(2003\)](#). The reduced Romerian game now satisfies the maintained conditions of Theorem 1 of [Frankel et al. \(2003\)](#): (a) supermodular payoffs in actions; (b) state-monotone payoffs; (c) continuous payoffs; (d) dominance regions on both ends. Theorem 1 of Frankel–Morris–Pauzner therefore applies: as $\sigma_\zeta \rightarrow 0$, iterated deletion of strictly dominated strategies converges to a unique strategy profile in monotone (threshold) form.

Step 6: Closed-form risk-dominant cutoff. At the cutoff $\phi_i = \phi^{rd}$, a Romerian must be indifferent between the two selections under Laplacian beliefs. The indifference condition reduces to

$$\int_0^1 [U_R(n_R^*(\bar{n}_q), \bar{n}_q) - U_R^0] dq = 0,$$

where $\bar{n}_q$ is the aggregate fertility implied by a fraction q of Romerians playing high. Equivalently,

$$\int_{\bar{n}^L}^{\bar{n}^H} u'_R(n_R^*(\bar{n}^e)) d\bar{n}^e = 2 \int_{\bar{n}^L}^{\bar{n}^{rd}} u'_R(n_R^*(\bar{n}^e)) d\bar{n}^e,$$

which under the log-quadratic specification has a unique closed-form solution $\bar{n}^{rd}$, and ϕ^{rd} is recovered from $T^{-1}(\bar{n}^{rd}) = \phi \bar{n}_R^{rd} + (1 - \phi) n_M^*(\bar{n}^{rd})$. The selected equilibrium is risk-dominant in the sense of [Frankel et al. \(2003\)](#). $\square$

A.5 Proof of Theorem 3

Step 1: Demographic-selection drift. For interior ϕ_t ,

$$\phi_t^{birth} - \phi_t = \frac{\phi_t(1 - \phi_t)(n_R - n_M)}{\bar{n}_t} > 0.$$

This drift is strictly positive at every interior ϕ_t and has no information dependence.

Step 2: Posterior dynamics under truth. Let $L_t = \ln[(1 - \pi_t)/\pi_t]$ be the log-odds in favor of Malthus. Bayes' rule (11) gives

$$L_{t+1} - L_t = \frac{(\theta_R - \theta_M) \ln(N_t/\bar{N})}{\sigma_\eta^2} \left[s_t - \bar{g} - \frac{\theta_R + \theta_M}{2} \ln(N_t/\bar{N}) \right],$$

which follows from expanding the log-likelihood ratio of two Gaussians. Under truth $\theta^* = \theta_R$, $\mathbb{E}^{\theta_R}[s_t \mid \mathcal{F}_t] = \bar{g} + \theta_R \ln(N_t/\bar{N})$, so

$$\mathbb{E}^{\theta_R}[L_{t+1} - L_t \mid \mathcal{F}_t] = -\mathcal{I}_t \leq 0.$$

Therefore L_t is a supermartingale under truth. By Doob's martingale convergence theorem, L_t converges a.s. to some $[-\infty, \infty]$ -valued limit L_∞ .

Step 3: $\sum_t \mathcal{I}_t = \infty$ almost surely. Suppose $\sum_t \mathcal{I}_t < \infty$ on a set of positive probability. Then $\mathcal{I}_t \rightarrow 0$, implying $\ln(N_t/\bar{N}) \rightarrow 0$, so $\bar{n}_t \rightarrow 1$ along a subsequence of full asymptotic density.

By Assumption 4, $\bar{n}_t \rightarrow 1$ requires $\phi_t \rightarrow \phi^{trap}$. On the same event, the cumulative posterior increments are summable, and through (13), the cumulative cultural-transmission drift accumulates without correction:

$$\phi_{t+1} - \phi_t \geq \lambda(\phi_t^{birth} - \phi_t) - (1 - \lambda) |\Delta\pi_{t+1}|,$$

with $\sum |\Delta\pi_{t+1}| < \infty$. Selection drift is bounded below by λc for some $c > 0$. Summing, ϕ_t increases beyond any bound below 1, contradicting $\phi_t \rightarrow \phi^{trap} < 1$. Hence $\sum_t \mathcal{I}_t = \infty$ a.s.

Step 4: $\pi_t \rightarrow 1$ almost surely. Combining $\mathbb{E}^{\theta_R}[\Delta L_t] = -\mathcal{I}_t$ with $\sum \mathcal{I}_t = \infty$, the law of large numbers for martingale differences gives $L_t \rightarrow -\infty$ a.s., hence $\pi_t \rightarrow 1$ a.s.

Step 5: $\phi_t \rightarrow 1$ almost surely. Given $\pi_t \rightarrow 1$, for any $\varepsilon > 0$ there exists $T(\omega)$ with $\pi_t > 1 - \varepsilon$ for $t \geq T$. The dynamics (13) give

$$1 - \phi_{t+1} \leq \lambda(1 - \phi_t) + (1 - \lambda)\varepsilon,$$

where the inequality uses $\phi_t^{birth} \geq \phi_t$. Iterating, $\limsup_{t \rightarrow \infty} (1 - \phi_t) \leq \varepsilon$. Since ε is arbitrary, $\phi_t \rightarrow 1$ a.s.

Step 6: $\bar{n}_t \rightarrow n_R^*$ and $\sup_t \ln(N_t/\bar{N}) = +\infty$. Convergence of fertility follows from continuity. Since $n_R^* > 1$, eventually $\bar{n}_t \geq (1 + n_R^*)/2 > 1$, so N_t grows geometrically. $\square$

A.6 Proof of Proposition 2

Throughout, the trajectory is assumed to remain in $\mathcal{T}(\delta, \zeta)$; the proof also verifies trap closure.

Part (i): Bound on $\mathbb{E}[\Delta\pi_t \mid \mathcal{F}_t]$. Using the chain rule on $\pi_t = (1 + e^{L_t})^{-1}$, $d\pi/dL = -\pi(1 - \pi)$. A second-order Taylor expansion:

$$\Delta\pi_t = -\pi_t(1 - \pi_t) \Delta L_t + \frac{1}{2}\pi_t(1 - \pi_t)(1 - 2\pi_t)(\Delta L_t)^2 + O((\Delta L_t)^3).$$

Taking expectations under truth and using $\mathbb{E}[(\Delta L_t)^2] = O(\mathcal{I}_t)$:

$$|\mathbb{E}[\Delta\pi_t \mid \mathcal{F}_t]| \leq \pi_t(1 - \pi_t) \mathcal{I}_t + O(\mathcal{I}_t).$$

In the trap region, $\mathcal{I}_t \leq \bar{\mathcal{I}}(\delta)$. Since $\pi(1 - \pi) \leq 1/4$, the leading term is bounded by $\bar{\mathcal{I}}(\delta)/4 = C_1 \bar{\mathcal{I}}(\delta)$ with $C_1 = 1/4$.

Part (ii): Bound on demographic-selection drift.

$$\phi_{t+1}^{birth} - \phi_t = \frac{\phi_t(1 - \phi_t)(n_R - n_M)}{\bar{n}_t} \leq \frac{(1 - \phi_t)(n_R - n_M)}{n_M},$$

where the inequality uses $\phi_t/\bar{n}_t \leq 1/n_M$. Therefore $C_2 = 1/n_M$.

Trap closure. If $(\phi_t, N_t, \pi_t) \in \mathcal{T}(\delta, \zeta)$, then $(\phi_{t+1}, N_{t+1}, \pi_{t+1}) \in \mathcal{T}(\delta, \zeta)$ for δ, ζ small enough. The condition $\bar{n}(\phi_{t+1}) \in [1, 1 + \zeta]$ follows from continuity. The condition $|\ln(N_{t+1}/\bar{N})| < \delta$ follows from $\bar{n}_t \in [1, 1 + \zeta]$. The condition $\pi_{t+1} \leq \phi_{t+1}$ follows from

$$\phi_{t+1} - \pi_{t+1} = \lambda(\phi_t^{birth} - \phi_t) + \lambda(\phi_t - \pi_t) - \lambda \Delta\pi_t,$$

where each summand is non-negative in expectation under truth in the trap.

Part (iii): Lower bound on $\mathbb{E}[T^{esc}]$. Combining (i) and (ii),

$$\mathbb{E}[\phi_{t+1} - \phi_t \mid \mathcal{F}_t] \leq \lambda C_2(n_R - n_M)(1 - \phi_t) + (1 - \lambda)C_1 \bar{\mathcal{I}}(\delta) \leq D.$$

Define $M_t = \phi_t - tD$. Then M_t is a supermartingale on the trap event. Optional stopping applied to the bounded stopping time $T \wedge n$:

$$D \mathbb{E}[T \wedge n] \geq \bar{\phi} \mathbb{P}(T \leq n) - \phi_0.$$

Letting $n \rightarrow \infty$,

$$\mathbb{E}[T^{esc}] \geq \frac{\bar{\phi} - \phi_0}{D}. \quad \square$$

A.7 Trap Duration (Auxiliary Result)

Corollary 3. *Under Proposition 2, when learning is the binding escape channel, $\mathbb{E}[T^{esc}] = \Omega(1/\delta^2)$ as $\delta \rightarrow 0$.*

Proof. $\bar{\mathcal{I}}(\delta) = O(\delta^2)$. With $D = a + b\delta^2$, we have $\mathbb{E}[T^{esc}] \geq (\bar{\phi} - \phi_0)/(a + b\delta^2)$. When the cultural channel is shut down or weak relative to the learning channel, a is small and $\mathbb{E}[T^{esc}] \geq \Omega(\delta^{-2})$. $\square$

A.8 Trap Boundary (Auxiliary Result)

Proposition 8. *The critical population at which the learning channel matches the demographic-selection channel in magnitude is*

$$\frac{N^{crit}(\phi)}{\bar{N}} = \exp \left(\sqrt{\frac{2\sigma_\eta^2 \cdot \lambda C_2(n_R - n_M)(1 - \phi)}{(1 - \lambda)C_1(\theta_R - \theta_M)^2}} \right).$$

Proof. Equate the per-period selection drift bound with the per-period learning drift and solve for $\ln(N/\bar{N})$. $\square$

A.9 Proof of Proposition 3

Part (i): Selection drift on births.

$$\phi_t^{birth} - \phi_t = \frac{\phi_t(1 - \phi_t)(n_R - n_M)}{\bar{n}_t}.$$

Since $\phi_t \in (0, 1)$ implies $\phi_t(1 - \phi_t) \in (0, 1/4]$, and $\bar{n}_t \geq n_M > 0$, the drift is strictly positive and bounded above by $(n_R - n_M)/(4n_M)$.

Part (ii): Learning drift on the posterior. Repeating the log-likelihood-ratio expansion under truth $\theta^* = \theta_M$:

$$L_{t+1} - L_t = \frac{(\theta_R - \theta_M) \ln(N_t/\bar{N})}{\sigma_\eta^2} \left[s_t - \bar{g} - \frac{\theta_R + \theta_M}{2} \ln(N_t/\bar{N}) \right].$$

Under $\theta^* = \theta_M$, $\mathbb{E}^{\theta_M}[s_t] = \bar{g} + \theta_M \ln(N_t/\bar{N})$. Substituting,

$$\mathbb{E}^{\theta_M}[L_{t+1} - L_t \mid \mathcal{F}_t] = +\mathcal{I}_t.$$

Using the chain-rule expansion $d\pi/dL = -\pi(1 - \pi)$,

$$\mathbb{E}^{\theta_M}[\pi_{t+1} - \pi_t \mid \mathcal{F}_t] = -\pi_t(1 - \pi_t) \mathcal{I}_t + o(\mathcal{I}_t).$$

This is strictly negative for $\pi_t \in (0, 1)$ and $\ln(N_t/\bar{N}) \neq 0$, and exactly zero at $\ln(N_t/\bar{N}) = 0$.

Part (iii): Sign change of the joint drift in ϕ_t . From $\phi_{t+1} = \lambda\phi_t^{birth} + (1 - \lambda)\pi_{t+1}$,

$$\phi_{t+1} - \phi_t = \lambda(\phi_t^{birth} - \phi_t) + (1 - \lambda)(\pi_t - \phi_t) + (1 - \lambda)(\pi_{t+1} - \pi_t).$$

Taking expectations under truth and using Part (ii):

$$\mathbb{E}^{\theta_M}[\phi_{t+1} - \phi_t \mid \mathcal{F}_t] = \underbrace{\lambda(\phi_t^{birth} - \phi_t)}_{>0} + \underbrace{(1 - \lambda)(\pi_t - \phi_t)}_{\text{sign of } \pi_t - \phi_t} + \underbrace{(1 - \lambda)[- \pi_t(1 - \pi_t)\mathcal{I}_t + o(\mathcal{I}_t)]}_{\leq 0}. \quad (35)$$

At $\ln(N_t/\bar{N}) = 0$: $\mathcal{I}_t = 0$, so the third term vanishes. The drift reduces to $\lambda(\phi_t^{birth} - \phi_t) + (1 - \lambda)(\pi_t - \phi_t)$, non-negative whenever $\pi_t \geq \phi_t - \frac{\lambda}{1 - \lambda}(\phi_t^{birth} - \phi_t)$.

For $|\ln(N_t/\bar{N})|$ large: the third term grows quadratically while the first two terms are bounded. The third term dominates and the drift becomes strictly negative once

$$\pi_t(1 - \pi_t) \mathcal{I}_t > \frac{\lambda(\phi_t^{birth} - \phi_t)}{1 - \lambda} + \max\{\pi_t - \phi_t, 0\}.$$

Substituting $\mathcal{I}_t$ and solving:

$$[\ln(N_t/\bar{N})]^2 > \bar{L}(\phi_t, \pi_t, \lambda)^2 \equiv \frac{2\sigma_\eta^2 [\lambda(\phi_t^{birth} - \phi_t)/(1 - \lambda) + \max\{\pi_t - \phi_t, 0\}]}{\pi_t(1 - \pi_t)(\theta_R - \theta_M)^2}.$$

Continuity implies the existence of a sign-changing state. □

A.10 State-Dependent Value of Announcement (Auxiliary Results)

Proposition 9. *Suppose the economy is in $\mathcal{T}(\delta, \zeta)$ in a Romerian world. Under W^{tot} ,*

$$\mathcal{V}^{trap}(\tau_\nu^{trap}) = P^{esc}(\tau_\nu^{trap}) \cdot \Delta W^{esc} - P^{deepen}(\tau_\nu^{trap}) \cdot \Delta W^{deepen},$$

with a threshold $\underline{\tau}^{trap}$ above which $\mathcal{V}^{trap} > 0$. Because ΔW^{esc} is large, $\underline{\tau}^{trap}$ is moderate.

Sketch of proof. An announcement of precision τ_ν shifts households' subjective θ^{post} toward the announced value, with magnitude proportional to $\tau_\nu/(\tau_\nu + \tau_{prior})$. In the trap, ϕ_t is below ϕ^{trap} . A truthful announcement raises subjective optimism, increasing n_R^* and n_M^* . If the resulting $\bar{n}_t$ exceeds 1, population grows, learning engages, and escape probability P^{esc} becomes positive. The deepening event has probability bounded above by $\mathbb{P}(\nu < -\nu_{min})$ for some threshold, which is small for moderate τ_ν . The threshold $\underline{\tau}^{trap}$ is moderate because ΔW^{esc} reflects the large welfare gain of escaping prolonged stagnation. $\square$

Proposition 10. *In a Malthusian world with non-monotone dynamics under W^{tot} ,*

$$\mathcal{V}^{cycle}(\tau_\nu) = \Omega(\text{Var}(\theta_i) - \mathbb{E}[\text{Var}(\theta_i^{post})]) - \Lambda((\bar{\theta} - \theta^*)^2 - \mathbb{E}[(\bar{\theta}^{post} - \theta^*)^2]) - C^{adj},$$

where $\Omega > 0$ is the marginal welfare gain from reducing dispersion, $\Lambda > 0$ from reducing mean bias, and $C^{adj} > 0$ is an adjustment cost. For sufficiently precise signals, $\mathcal{V}^{cycle} > 0$.

Sketch of proof. The decomposition (28) provides the mapping from variance and bias of beliefs to welfare losses Δ_3 and Δ_2 . Both variance and bias are decreasing in τ_ν . The adjustment cost reflects the welfare loss from disrupting agents' contingent plans during a cycle. $\square$

B Welfare Analysis Proofs and Auxiliary Results

B.1 Proof of Proposition 4

Step 1: Planner's Bellman. Define

$$V^P(A_t, N_t) = \max_{n_t} \left\{ N_t [\ln A_t + \ln(1 - \psi(n_t))] + \beta \mathbb{E}^*[V^P(A_{t+1}, n_t N_t)] \right\}.$$

Conjecture $V^P(A, N) = N \cdot w^P(A, N)$ with $w^P(A, N) \approx v^P + \xi_A^P \ln A + \xi_N^P \ln(N/\bar{N})$. Substituting and dividing by N ,

$$w^P(A, N) = \ln A + \ln(1 - \psi(n)) + \beta n \mathbb{E}^*[w^P(A', N')].$$

The planner's per-capita Bellman has the same form as the household's, but the planner's FOC additionally internalizes the population-multiplier through differentiating $N_{t+1} = nN_t$.

Step 2: Coefficient matching.

$$\xi_A^P = \frac{1}{1 - \beta n^{FB}}, \quad \xi_N^P = \frac{\beta n^{FB} \theta^*}{(1 - \beta n^{FB})^2}.$$

Step 3: Planner FOC. Differentiating with respect to n_t and noting $N_{t+1} = n_t N_t$, at the reference $\partial w^P / \partial N = \xi_N^P / N$, so the marginal value of population is $V_N^P = w^P + \xi_N^P$. The FOC reads (27).

Step 4: Comparison with household. The household's common-correct-belief FOC is $\gamma n^{CB} = \beta \mathbb{E}[V_{CB}(A', N')]$ where $V_{CB}(A, N) = w^P(A, N)$. The wedge between (27) and the household FOC is exactly $\beta \xi_N^P$, with the sign of θ^* . Strict monotonicity of γn gives:

$$\theta^* > 0 \implies n^{FB} > n^{CB}, \quad \theta^* < 0 \implies n^{FB} < n^{CB}.$$

Step 5: $n^{CB} > n_M^*$. By Lemma 1, fertility is increasing in θ_j . With $\theta^* = \theta_R > 0$, $n^{CB} = n^*(\theta_R) > n_M^* = n^*(\theta_M)$. Combined with Step 4, $n^{FB} > n^{CB} > n_M^*$.

Step 6: Welfare decomposition. The decomposition (28) is established below. □

B.2 First-Best under Average Utilitarianism (Auxiliary Result)

Proposition 11. *In a Romerian world, $n^{CB} < n^{FB, avg} < n^{FB, tot}$. The average-utilitarian planner internalizes the per-capita scale-effect channel through $\ln A_t$, but lacks the population-multiplier term N_t that gives the total utilitarian planner the additional wedge.*

Proof. The average-utilitarian Bellman is $V^{P, avg}(A, N) = \ln A + \ln(1 - \psi(n)) + \beta \mathbb{E}^*[V^{P, avg}(A', N')]$, without the multiplicative N . Matching gives $\xi_N^{P, avg} = \beta n^{FB, avg} \theta^* / (1 - \beta n^{FB, avg})^2$. The wedge under total utilitarianism is $\beta(w^P + \xi_N^P)$ and under average utilitarianism is $\beta \xi_N^{P, avg}$. With $w^P > 0$, $\beta \xi_N^P + \beta w^P > \beta \xi_N^{P, avg} > 0$. Strict monotonicity of γn gives $n^{FB, tot} > n^{FB, avg} > n^{CB}$. □

B.3 Proof of Proposition 5

Decompose by inserting two intermediate allocations:

$$W^{FB}(\theta^*) - W^{HB} = \underbrace{[W^{FB}(\theta^*) - W^{CB}(\theta^*)]}_{\Delta_1} + \underbrace{[W^{CB}(\theta^*) - W^{CB}(\bar{\theta})]}_{\Delta_2} + \underbrace{[W^{CB}(\bar{\theta}) - W^{HB}]}_{\Delta_3},$$

where $W^{CB}(\theta)$ is the welfare under decentralized common-belief equilibrium when households all believe θ ; W^{HB} is the heterogeneous-beliefs welfare.

Δ_1 (externality). $\Delta_1 \geq 0$, with equality iff $\theta^* = 0$. At a quadratic approximation around n^{CB} ,

$$\Delta_1 \approx \frac{1}{2}\omega(\beta\xi_N^P)^2/\gamma,$$

where $\omega > 0$ is the curvature of welfare in fertility.

Δ_2 (mean bias). Where $\bar{\theta} = \phi\theta_R + (1 - \phi)\theta_M$,

$$\Delta_2 \approx \frac{1}{2}\omega(n^{CB}(\bar{\theta}) - n^{CB}(\theta^*))^2 \approx \frac{1}{2}\omega\kappa^2(\bar{\theta} - \theta^*)^2,$$

where $\kappa = \partial n^{CB}/\partial\theta$.

Δ_3 (dispersion). Heterogeneous beliefs induce dispersion $\text{Var}(n_i) \approx \kappa^2\text{Var}(\theta_i)$, so

$$\Delta_3 \approx \frac{1}{2}\omega\kappa^2\text{Var}(\theta_i).$$

The expansion is around n^{SP} , so Δ_1 explicitly carries the externality.

Sign and magnitude properties. $\Delta_1, \Delta_2, \Delta_3 \geq 0$ by construction. Δ_2 increases in $(\bar{\theta} - \theta^*)^2$; Δ_3 increases in $\text{Var}(\theta_i)$. Public announcements affect $\bar{\theta}^{post}$ and $\text{Var}(\theta_i^{post})$ but not Δ_1 . $\square$

B.4 Proof of Theorem 4

Under W^{dyn} ,

$$W^{dyn}(\mu) = \phi_0 V_R^{post}(\mu) + (1 - \phi_0) V_M^{post}(\mu).$$

By Blackwell's theorem applied to the household decision problem under subjective preferences, any signal that is more informative weakly increases V_j^{post} . Truthful announcement

$\mu = \hat{\theta}$ is most informative among signals with prior mean $\hat{\theta}$. There is no externality the planner cares about beyond what households themselves would: W^{dyn} uses subjective beliefs, so the externality wedge does not enter. The optimal announcement is $\mu^{dyn} = \hat{\theta}$. $\square$

B.5 Proof of Theorem 5

Under W^{tot} , the announcement μ shifts households' subjective θ^{post} toward μ by a factor $\alpha = \tau_A/(\tau_A + \tau_{prior}) \in (0, 1)$. The pass-through affects subjective fertility approximately linearly:

$$n_j^{post} \approx n^{CB}(\theta_j) + \alpha\kappa(\mu - \theta_j).$$

Aggregate fertility:

$$\bar{n}^{post}(\mu) \approx \bar{n}^{HB} + \alpha\kappa(\mu - \bar{\theta}).$$

The total utilitarian welfare evaluated under truth, to second order around n^{FB} ,

$$W^{tot}(\bar{n}^{post}) \approx W^{FB} - \frac{1}{2}\Omega(\bar{n}^{post} - n^{FB})^2 - \frac{1}{2}\Omega'\text{Var}(n_i^{post}).$$

Differentiating with respect to μ :

$$\frac{dW^{tot}}{d\mu} = -\Omega(\bar{n}^{post} - n^{FB}) \cdot \alpha\kappa + \text{terms from } \text{Var}(n_i^{post}).$$

Setting equal to zero:

$$\bar{n}^{post,*} = n^{FB} \iff \mu^{tot} = \bar{\theta} + \frac{n^{FB} - \bar{n}^{HB}}{\alpha\kappa}.$$

Decomposing $n^{FB} - \bar{n}^{HB} = (n^{FB} - n^{CB}(\theta^*)) + (n^{CB}(\theta^*) - \bar{n}^{HB})$, the first term is the externality wedge and the second the bias due to heterogeneous beliefs. Approximating $n^{CB}(\theta^*) - \bar{n}^{HB} \approx \kappa(\theta^* - \bar{\theta})$:

$$\mu^{tot} \approx \theta^* + \frac{\beta\xi_N^P}{\alpha\kappa\gamma}.$$

Replacing θ^* by $\hat{\theta}$:

$$\mu^{tot} = \hat{\theta} + \chi^{tot}, \quad \chi^{tot} = \frac{\beta\xi_N^P}{\alpha\kappa\gamma}.$$

The shade $\chi^{tot} > 0$ iff $\theta^* > 0$.

Comparative statics. (i) ξ_N^P is increasing in $|\theta^*|$, so $|\chi^{tot}|$ is increasing in $|\theta^*|$. (ii) Evaluating ξ_N^P at non-zero deviation of population from reference involves a correction proportional to $|\ln(N_0/\bar{N})|$. (iii) Larger $\text{Var}(\theta_{i,0})$ raises curvature Ω' and lowers α , reducing χ^{tot} . (iv)

Higher τ_A raises α ; the absolute welfare effect of the announcement is increasing in τ_A .¹⁹ $\square$

B.6 Proof of Corollary 1

Repeating the argument of Theorem 5 with W^{avg} in place of W^{tot} , the only change is that the externality wedge becomes $\beta\xi_{N,avg}^P$ instead of $\beta\xi_N^P$. By Proposition 11, $0 < \xi_{N,avg}^P < \xi_N^P$ when $\theta^* > 0$. Therefore $0 < \chi^{avg} < \chi^{tot}$. $\square$

B.7 Proof of Corollary 2

Direct from Theorem 4, Corollary 1, and Theorem 5, with $0 < \chi^{avg} < \chi^{tot}$ in a Romerian world. $\square$

B.8 Proof of Proposition 6

The proof proceeds in three steps.

Step 1: Private validation. Suppose a single Romerian dynasty's prior is shifted from θ_R to $\theta_R + q$, leaving all other priors unchanged. The dynasty's value function depends on its own belief through the FOC and on the equilibrium aggregate $\bar{n}$, which depends on the priors of all other agents but not on this single dynasty's prior. Because the prior shift is private, $\bar{n}$ is unaffected. The marginal effect on V_R is

$$\left. \frac{dV_R^{dyn}}{dq} \right|_{private} = \beta\xi_{A,R} \ln(N/\bar{N}) + (\beta n_R / (1 - \beta n_R^*)) \ln \bar{n} + O(\text{small}),$$

which I normalize as the unit private benefit b^{priv} .

Step 2: Public validation. Suppose every household's prior is shifted by z . The household's own marginal effect is the same as in private validation, but the equilibrium aggregate $\bar{n}$ now responds: by Proposition 1 and equation (18), in a fully homogenized Romerian configuration,

$$\left. \frac{d\bar{n}}{dz} \right|_{public} = \rho(\phi).$$

Differentiating V_R with respect to z :

$$\left. \frac{dV_R^{dyn}}{dz} \right|_{public} = b^{priv} + \xi_{N,R} \rho(\phi) \cdot (1/\bar{n}).$$

¹⁹The original theorem statement is best interpreted as the absolute welfare effect of the announcement (which is increasing in τ_A), not the magnitude of the shading.

Under the bounded-range parameterization with half-width a , this can be written equivalently as

$$\left. \frac{dV_R^{dyn}}{dz} \right|_{public} = \frac{2a}{2a + \Delta_{\text{wedge}}} \cdot b^{priv},$$

where $\Delta_{\text{wedge}} = \xi_{N,R} \rho(\phi) / (\bar{n} \cdot b^{priv})$. The exact form under the uniform belief-dispersion family is $\Delta_{\text{wedge}} = \beta^2 n_R^* \theta_R / [\gamma \bar{n} (1 - \beta n_R^*)^2] = d_R$.

Step 3: Properties of the retained fraction. $2a/(2a + \Delta_{\text{wedge}}) \in (0, 1)$ because $a > 0$ and $\Delta_{\text{wedge}} > 0$ in a Romerian world. Differentiating with respect to Δ_{wedge} :

$$\frac{\partial}{\partial \Delta_{\text{wedge}}} \left(\frac{2a}{2a + \Delta_{\text{wedge}}} \right) = -\frac{2a}{(2a + \Delta_{\text{wedge}})^2} < 0,$$

so the retained fraction is strictly decreasing in the strategic-complementarity wedge.

When $\rho(\phi)$ is small (as in the trap), Δ_{wedge} is small and the retained fraction is close to one. In the multiplicity region where $\rho(\phi) > 1$, the wedge is larger and the retained fraction is smaller. $\square$

B.9 Proof of Proposition 7

In the multiplicity region, the welfare derivative with respect to μ includes a term reflecting the change in the realized selection at the cutoff ϕ^{rd} . Let $g(\phi)$ denote the prior density of ϕ . The welfare gain from a marginal lowering of the cutoff in a Romerian world is

$$\frac{dW^{tot}}{d\phi^{rd}} = -g(\phi^{rd}) \cdot [W^{tot}(\bar{n}^H, \cdot) - W^{tot}(\bar{n}^L, \cdot)],$$

which is negative because the bracket is positive. The chain rule gives

$$\left. \frac{dW^{tot}}{d\mu} \right|_{\text{selection}} = \frac{dW^{tot}}{d\phi^{rd}} \cdot \frac{d\phi^{rd}}{d\mu},$$

where $d\phi^{rd}/d\mu < 0$ in a Romerian world. The product is positive, so the selection-shifting marginal benefit is positive.

Adding this to the FOC of Theorem 5 yields an additional shading term χ^{sel} . The non-monotonicity follows from $g(\phi^{rd})$: when ϕ is far from ϕ^{rd} on either side with concentrated prior, $g(\phi^{rd})$ is small. χ^{sel} peaks when the prior places maximum mass near the cutoff. $\square$